

Chicago_3

August 27, 2024

1 Data Management and Predictions

Let's create an advanced aggregated dataset and then use it for predictive analysis. We'll do this in two steps: first, we'll aggregate the data and save it as a CSV file, and then we'll use that file for our predictive analysis.

1.1 1. Aggregation

This script creates a panel data by the following steps:

1. Loads the original crime data.
2. Creates a function to aggregate the data, counting total crimes, total arrests, and individual counts for each crime type and location description.
3. Aggregates the data by District and Date.
4. Saves the aggregated data to a new CSV file.

```
[1]: import pandas as pd
import numpy as np
```

```
[10]: # Load the original data
file_path = '/Users/YigitAydede/Library/CloudStorage/Dropbox/Documents/Courses/
↳MBAN/NLPBootcamp/PythonBC/Crimes_-_2024_20240804.csv'
df = pd.read_csv(file_path, parse_dates=['Date'])

# Convert Date to datetime
df['Date'] = pd.to_datetime(df['Date'], format='%m/%d/%Y %I:%M:%S %p',
↳errors='coerce')
```

```
/var/folders/b2/gpnsjh9j6bv5prtx7w5lsym80000gp/T/ipykernel_85973/3313344229.py:3
: UserWarning: Could not infer format, so each element will be parsed
individually, falling back to `dateutil`. To ensure parsing is consistent and
as-expected, please specify a format.
df = pd.read_csv(file_path, parse_dates=['Date'])
```

```
[13]: def aggregate_data(group):
# Initialize the dictionary with total crime and total arrest
agg = {
'Total_Crime': group['ID'].count(),
'Total_Arrest': group['Arrest'].sum()
}
```

```

    }

    # Add counts for each crime type
    for crime_type in df['Primary Type'].unique():
        agg[f'Crime_{crime_type.replace(" ", "_")}'] = (group['Primary Type']_
        == crime_type).sum()

    # Handle NaN values in 'Location Description' and add counts for each_
    location description
    group['Location Description'] = group['Location Description'].
    fillna('Unknown')
    for location in df['Location Description'].unique():
        agg[f'Location_{location.replace(" ", "_")}'] = (group['Location_
        Description'] == location).sum()

    return pd.Series(agg)

```

Ask Copilot for help to see what this code does.

```

[14]: # Aggregate data by District and Date
aggregated_df = df.groupby(['District', df['Date'].dt.date]).
    apply(aggregate_data)

```

/var/folders/b2/gpnsjh9j6bv5prtx7w5lsym80000gp/T/ipykernel_85973/1280444553.py:2
: DeprecationWarning: DataFrameGroupBy.apply operated on the grouping columns.
This behavior is deprecated, and in a future version of pandas the grouping
columns will be excluded from the operation. Either pass `include_groups=False`
to exclude the groupings or explicitly select the grouping columns after groupby
to silence this warning.

```

    aggregated_df = df.groupby(['District',
df['Date'].dt.date]).apply(aggregate_data)

```

```

[15]: print(aggregated_df.head())

```

		Total_Crime	Total_Arrest	Crime_MOTOR_VEHICLE_THEFT	\
District	Date				
1	2024-01-01	34	9		4
	2024-01-02	29	11		2
	2024-01-03	39	3		3
	2024-01-04	34	7		0
	2024-01-05	41	8		0

		Crime_DECEPTIVE_PRACTICE	Crime_THEFT	Crime_BATTERY	\
District	Date				
1	2024-01-01		1	8	12
	2024-01-02		2	9	5
	2024-01-03		6	18	3
	2024-01-04		4	11	7

2024-01-05	7	15	4
------------	---	----	---

District	Date	Crime_SEX_OFFENSE	Crime_CRIMINAL_DAMAGE	\
1	2024-01-01	0	0	
	2024-01-02	0	2	
	2024-01-03	0	1	
	2024-01-04	0	1	
	2024-01-05	0	1	

District	Date	Crime_CRIMINAL_SEXUAL_ASSAULT	\
1	2024-01-01	0	
	2024-01-02	0	
	2024-01-03	0	
	2024-01-04	0	
	2024-01-05	0	

District	Date	Crime_OFFENSE_INVOLVING_CHILDREN	...	\
1	2024-01-01	0	...	
	2024-01-02	0	...	
	2024-01-03	0	...	
	2024-01-04	0	...	
	2024-01-05	0	...	

District	Date	Location_CREDIT_UNION	\
1	2024-01-01	0	
	2024-01-02	0	
	2024-01-03	0	
	2024-01-04	0	
	2024-01-05	0	

District	Date	Location_VEHICLE_-_COMMERCIAL:_TROLLEY_BUS	\
1	2024-01-01	0	
	2024-01-02	0	
	2024-01-03	0	
	2024-01-04	0	
	2024-01-05	0	

District	Date	Location_HALLWAY	Location_GARAGE	Location_OFFICE	\
1	2024-01-01	0	0	0	
	2024-01-02	0	0	0	
	2024-01-03	0	0	0	
	2024-01-04	0	0	0	

District	Date	Location_RETAIL_STORE	Location_LIQUOR_STORE	\
1	2024-01-01	0	0	
	2024-01-02	0	0	
	2024-01-03	0	0	
	2024-01-04	0	0	
	2024-01-05	0	0	

District	Date	Location_CHA_HALLWAY	Location_CTA_"L"_TRAIN	\
1	2024-01-01	0	0	
	2024-01-02	0	0	
	2024-01-03	0	0	
	2024-01-04	0	0	
	2024-01-05	0	0	

District	Date	Location_CTA_TRACKS_-_RIGHT_OF_WAY
1	2024-01-01	0
	2024-01-02	0
	2024-01-03	0
	2024-01-04	1
	2024-01-05	0

[5 rows x 156 columns]

```
[16]: # Write the aggregated DataFrame to a CSV file
output_file_path = 'aggregated_crime_data.csv' # Replace with your desired
        ↳ output file path
aggregated_df.to_csv(output_file_path)
```

1.2 2. Chicago Crime Predictive Analysis

This script performs the following analysis:

1. Calculates and visualizes a correlation matrix for all crime types.
2. Defines a function to predict one crime type based on another using linear regression.
3. Finds the most correlated crime types and attempts to predict one based on the other.
4. As an example, it also predicts ASSAULT based on THEFT.

```
[17]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
from sklearn.metrics import mean_squared_error, r2_score
```


```
[18]: # Load the aggregated data
file_path = '/Users/YigitAydede/Library/CloudStorage/Dropbox/Documents/Courses/
↳MBAN/NLPBootcamp/PythonBC/aggregated_crime_data.csv'
df = pd.read_csv(file_path)
```

1.2.1 2.1 Correlation

```
[19]: # Select crime type columns (assuming they start with 'Crime_')
crime_columns = [col for col in df.columns if col.startswith('Crime_')]

# Calculate correlation matrix for crime types
correlation_matrix = df[crime_columns].corr()

# Visualize the correlation matrix
plt.figure(figsize=(20, 16))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix of Crime Types')
plt.tight_layout()
plt.show()
```



```
[23]: # Select crime type columns (assuming they start with 'Crime_')
crime_columns = [col for col in df.columns if col.startswith('Crime_')]

# Calculate correlation matrix for crime types
correlation_matrix = df[crime_columns].corr()

# Find pairs with correlation greater than 0.2
correlation_threshold = 0.2
high_correlation_pairs = []
for i in range(len(crime_columns)):
    for j in range(i+1, len(crime_columns)):
        if correlation_matrix.iloc[i, j] > correlation_threshold:
            high_correlation_pairs.append((crime_columns[i], crime_columns[j],
            correlation_matrix.iloc[i, j]))

# Print the high correlation pairs
for pair in high_correlation_pairs:
```

```
print(f"Pair: {pair[0]} - {pair[1]}, Correlation: {pair[2]}")
```

```
Pair: Crime_MOTOR_VEHICLE_THEFT - Crime_CRIMINAL_DAMAGE, Correlation:
0.24962841130447888
Pair: Crime_DECEPTIVE_PRACTICE - Crime_THEFT, Correlation: 0.35668458424264426
Pair: Crime_THEFT - Crime_CRIMINAL_TRESPASS, Correlation: 0.23231920724746682
Pair: Crime_BATTERY - Crime_CRIMINAL_DAMAGE, Correlation: 0.2531180711829242
Pair: Crime_BATTERY - Crime_ASSAULT, Correlation: 0.31328139117196047
Pair: Crime_BATTERY - Crime_WEAPONS_VIOLATION, Correlation: 0.27826439102933564
Pair: Crime_CRIMINAL_DAMAGE - Crime_ASSAULT, Correlation: 0.2336953556787837
Pair: Crime_OTHER_OFFENSE - Crime_ASSAULT, Correlation: 0.2020758494835316
Pair: Crime_ASSAULT - Crime_WEAPONS_VIOLATION, Correlation: 0.23104828841621122
```

```
[26]: # Find the most correlated crime types
top_correlations = correlation_matrix.unstack().sort_values(ascending=False)
top_correlations = top_correlations[(top_correlations > 0.2)]
```

1.2.2 2.2 Predicting Theft

To create a predictive model that predicts one type of crime (e.g., theft) using all other crime types and location descriptions as features, you can follow these steps:

Let's have only the theft data. We predict by using its past 3 days within the district. Then we use more complex data to predict the next day.

```
[30]: # Load the aggregated data
file_path = '/Users/YigitAydede/Library/CloudStorage/Dropbox/Documents/Courses/
↳MBAN/NLPBootcamp/PythonBC/aggregated_crime_data.csv'
aggregated_df = pd.read_csv(file_path, parse_dates=['Date'])

# Ensure the data is sorted by date for each district
aggregated_df.sort_values(by=['District', 'Date'], inplace=True)

# Inspect the data
print(aggregated_df.head())
```

	District	Date	Total_Crime	Total_Arrest	Crime_MOTOR_VEHICLE_THEFT	\
0	1	2024-01-01	34	9	4	
1	1	2024-01-02	29	11	2	
2	1	2024-01-03	39	3	3	
3	1	2024-01-04	34	7	0	
4	1	2024-01-05	41	8	0	

	Crime_DECEPTIVE_PRACTICE	Crime_THEFT	Crime_BATTERY	Crime_SEX_OFFENSE	\
0	1	8	12	0	
1	2	9	5	0	
2	6	18	3	0	
3	4	11	7	0	
4	7	15	4	0	

	Crime_CRIMINAL_DAMAGE	...	Location_CREDIT_UNION	\
0	0	...	0	
1	2	...	0	
2	1	...	0	
3	1	...	0	
4	1	...	0	

	Location_VEHICLE_-_COMMERCIAL:_TROLLEY_BUS	Location_HALLWAY	\
0	0	0	
1	0	0	
2	0	0	
3	0	0	
4	0	0	

	Location_GARAGE	Location_OFFICE	Location_RETAIL_STORE	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	Location_LIQUOR_STORE	Location_CHA_HALLWAY	Location_CTA_"L"_TRAIN	\
0	0	0	0	
1	0	0	0	
2	0	0	0	
3	0	0	0	
4	0	0	0	

	Location_CTA_TRACKS_-_RIGHT_OF_WAY
0	0
1	0
2	0
3	1
4	0

[5 rows x 158 columns]

Create lag features

```
[31]: def create_lag_features(df, target, lags, groupby_col):
    for lag in range(1, lags + 1):
        df[f'{target}_lag_{lag}'] = df.groupby(groupby_col)[target].shift(lag)
    return df

# Define the number of lag days
num_lags = 3
```



```

# Create lag features for the 'Crime_THEFT' column
aggregated_df = create_lag_features(aggregated_df, 'Crime_THEFT', num_lags,
↪ 'District')

# Drop rows with NaN values that result from lag creation
aggregated_df.dropna(inplace=True)

# Inspect the data with lag features
print(aggregated_df.head())

```

	District	Date	Total_Crime	Total_Arrest	Crime_MOTOR_VEHICLE_THEFT	\
3	1	2024-01-04	34	7	0	
4	1	2024-01-05	41	8	0	
5	1	2024-01-06	30	6	2	
6	1	2024-01-07	22	5	3	
7	1	2024-01-08	34	8	3	

	Crime_DECEPTIVE_PRACTICE	Crime_THEFT	Crime_BATTERY	Crime_SEX_OFFENSE	\
3	4	11	7	0	
4	7	15	4	0	
5	4	9	2	0	
6	3	7	2	0	
7	1	10	4	0	

	Crime_CRIMINAL_DAMAGE	...	Location_GARAGE	Location_OFFICE	\
3	1	...	0	0	
4	1	...	0	0	
5	5	...	0	0	
6	4	...	0	0	
7	1	...	0	0	

	Location_RETAIL_STORE	Location_LIQUOR_STORE	Location_CHA_HALLWAY	\
3	0	0	0	
4	0	0	0	
5	0	0	0	
6	0	0	0	
7	0	0	0	

	Location_CTA_"L"_TRAIN	Location_CTA_TRACKS_-_RIGHT_OF_WAY	\
3	0	1	
4	0	0	
5	0	0	
6	0	0	
7	0	0	

	Crime_THEFT_lag_1	Crime_THEFT_lag_2	Crime_THEFT_lag_3
3	18.0	9.0	8.0
4	11.0	18.0	9.0

5	15.0	11.0	18.0
6	9.0	15.0	11.0
7	7.0	9.0	15.0

[5 rows x 161 columns]

```
[32]: # Define the target variable
target = 'Crime_THEFT'

# Define the feature set by excluding the target variable and other non-feature
↳ columns
feature_cols = [f'Crime_THEFT_lag_{i}' for i in range(1, num_lags + 1)]

# Select the features and target from the DataFrame
X = aggregated_df[feature_cols]
y = aggregated_df[target]
```

Split the data into training and testing sets

```
[33]: from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

A Simple linear regression model

```
[34]: from sklearn.linear_model import LinearRegression

# Initialize the model
model = LinearRegression()

# Train the model
model.fit(X_train, y_train)
```

```
[34]: LinearRegression()
```

Evaluate the model

```
[35]: from sklearn.metrics import mean_squared_error, r2_score

# Make predictions on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"Mean Squared Error: {mse}")
print(f"R^2 Score: {r2}")
```

Mean Squared Error: 11.396217555270395

R² Score: 0.46141715844335973

Make predictions

```
[37]: # Make predictions on new data (e.g., the test set)
      predictions = model.predict(X_test)

      # Create a DataFrame with the actual and predicted values along with the dates
      ↪and districts
      result_df = X_test.copy()
      result_df['Actual_Crime_THEFT'] = y_test
      result_df['Predicted_Crime_THEFT'] = predictions
      result_df['District'] = aggregated_df.loc[X_test.index, 'District']
      result_df['Date'] = aggregated_df.loc[X_test.index, 'Date']

      # Display the results
      print(result_df.head())
```

	Crime_THEFT_lag_1	Crime_THEFT_lag_2	Crime_THEFT_lag_3	\
2438	9.0	14.0	7.0	
1940	2.0	4.0	3.0	
975	7.0	2.0	0.0	
2384	9.0	10.0	13.0	
4147	5.0	4.0	4.0	

	Actual_Crime_THEFT	Predicted_Crime_THEFT	District	Date
2438	8	9.370871	12	2024-05-22
1940	7	3.808013	10	2024-03-02
975	6	4.077546	5	2024-05-19
2384	14	9.867449	12	2024-03-29
4147	13	4.958095	22	2024-06-30