

Python_for_data

January 6, 2024

1 Notebook for Data Analytics in Python

Python was not written with data analysis in mind. It's a general purpose programming language!

This notebook is similar to Rmarkdown, but less practical:-) Knitting a Jupyter Notebook to a PDF is a multi-step process that typically involves converting the notebook to a different format first (like HTML or LaTeX) and then to a PDF. Here's a general process you can follow:

- `nbconvert` is a tool provided by Jupyter that allows you to convert notebooks to various other formats, including HTML and PDF. First, ensure that `nbconvert` is installed. Run this command in your terminal:`conda install nbconvert`
- To convert a notebook to PDF, you need a TeX distribution installed. `nbconvert` uses TeX to generate PDFs. For Windows, you can use MikTeX. For macOS, MacTeX is a common choice.
- `jupyter nbconvert --to pdf YourNotebook.ipynb`
- `jupyter nbconvert --to html YourNotebook.ipynb`

2 Introduction

The following command instructs Python to join together the numbers 3, 4, and 5, and to save them as a list named `x`. When we type `x`, it gives us back the list.

```
[1]: x = [3, 4, 5]
x
```

```
[1]: [3, 4, 5]
```

Note that we used the brackets `[]` to construct this list. We will often want to add two sets of numbers together. It is reasonable to try the following code, though it will not produce the desired results.

```
[2]: y = [4, 9, 7]
x + y
```

```
[2]: [3, 4, 5, 4, 9, 7]
```

The result may appear slightly counterintuitive: why did Python not add the entries of the lists element-by-element? In Python, lists hold arbitrary objects, and are added using concatenation.

In fact, concatenation is the behavior that we saw earlier when we entered “hello” + ” ” + “world”. This example reflects the fact that Python is a general-purpose programming language. Much of Python’s data-specific functionality comes from other packages, notably `numpy` and `pandas`. We’ll see a lot but here is the starter:

```
[1]: import numpy as np
      x = np.array([3, 4, 5])
      y = np.array([4, 9, 7])
      x + y
```

```
[1]: array([ 7, 13, 12])
```

In `numpy`, matrices are typically represented as two-dimensional arrays, and vectors as one-dimensional arrays. We’ll see `numpy` later in Chapter 9 more detail. Let’s start with simple examples without calling a library

```
[ ]: Create two variables named a and b so that they contain the following
      strings respectively: "23 to 0" and "C'est la piquette, Jack!".
```

```
[1]: a = "23 to 0"
      b = "C'est la piquette, Jack!"
```

```
[2]: print(a)
      print(b)
```

23 to 0

C'est la piquette, Jack!

Display the number of characters from `a`, then `b`.

```
[25]: print(len(a))
      print(len(b))
```

7

24

Concatenate `a` and `b` in a single string, adding a comma as a separating character.

```
[29]: print(a + ", " + b)
```

23 to 0, C'est la piquette, Jack!

Same question by choosing a separation that allows a line break between the two sentences.

```
[32]: print(a + "\n" + b)
```

23 to 0

C'est la piquette, Jack!

Using the appropriate method, capitalize `a` and `b`.

```
[42]: a = a.upper()
      b = b.upper()
      print(a)
      print(b)
```

```
23 TO 0
C'EST LA PIQUETTE, JACK!
```

Using the appropriate method, lowercase a and b.

```
[44]: a = a.lower()
      b = b.lower()
      print(a)
      print(b)
```

```
23 to 0
c'est la piquette, jack!
```

Extract the word la and Jack from the string b, using indexes.

```
[58]: #la
      ind = b.find("la")
      l = len("la")
      print(b[ind:ind + l])
      #Jack
      ind = b.find("jack")
      l = len("jack")
      print(b[ind:ind + l])
```

```
la
jack
```

Look for the sub-chain piqu in b, then do the same with the sub-chain mauvais.

3 Types of Data

str(), int(), float(), date()

```
[63]: x = 1.5
      print(type(x))
      print(int(x))
      print(x == 1)
```

```
<class 'float'>
1
False
```

4 Structures

4.1 Lists

```
[64]: x = ["Pascaline", "Gauthier", "Xuan", "Jimmy"]
      print(x)
```

```
['Pascaline', 'Gauthier', 'Xuan', 'Jimmy']
```

```
[65]: z = ["Piketty", "Thomas", 1971]
      print(z)
```

```
['Piketty', 'Thomas', 1971]
```

```
[67]: print(x[0]) # The first element of x
      print(x[-1]) # The last element of x
```

```
Pascaline
Jimmy
```

```
[68]: print(x[1:2]) # The first and second elements of x
      print(x[2:]) # From the second element (not included) to the end of x
      print(x[:-2]) # From the first to the penultimate (not included)
```

```
['Gauthier']
['Xuan', 'Jimmy']
['Pascaline', 'Gauthier']
```

```
[77]: tweets = ["aaa", "bbb"]
      followers = ["Anne", "Bob", "Irma", "John"]
      conuts = [tweets, followers]
      res = conuts[1][3] # The 4th item of the 2nd item on the list counts
      print(res)
```

```
John
```

```
[80]: print(len(conuts))
      print(len(conuts[1]))
```

```
2
4
```

4.1.1 Modifications

```
[81]: x = [1, 3, 5, 6, 9]
      x[3] = 7 # Replacing the 4th element
      print(x)
```

```
[1, 3, 5, 7, 9]
```

```
[83]: x.append(11) # Add value 11 at the end of the list
      print(x)
      y = [13, 15]
      x.extend(y)
      print(x)
```

```
[1, 3, 5, 7, 9, 11, 11]
[1, 3, 5, 7, 9, 11, 11, 13, 15]
```

```
[84]: x.remove(3) # Remove the fourth element
      print(x)
      x = [1, 3, 5, 6, 9]
      del x[3] # Remove the fourth element
      print(x)
```

```
[1, 5, 7, 9, 11, 11, 13, 15]
```

```
[85]: x = [1, 3, 5, 6, 10]
      x[3:5] = [7, 9] # Replaces 4th and 5th values
      print(x)
```

```
[1, 3, 5, 7, 9]
```

```
[86]: x = [1, 2, 3, 4, 5]
      x[2:3] = ['a', 'b', 'c', 'd'] # Replaces the 3rd value
      print(x)
```

```
[1, 2, 'a', 'b', 'c', 'd', 4, 5]
```

Verifying if a Value is Present

```
[87]: x = [1, 2, 3, 4, 5]
      print(1 in x)
```

True

Be careful, copying a list is not trivial in Python. Let's take an example.

```
[88]: x = [1, 2, 3]
      y = x
      y[0] = 0
      print(y)
      print(x)
```

```
[0, 2, 3]
[0, 2, 3]
```

To copy a list, there are several ways to do so. Among them, the use of the `list()` function:

```
[89]: x = [1, 2, 3]
      y = list(x)
      y[0] = 0
      print("x : ", x)
      print("y : ", y)
```

```
x : [1, 2, 3]
y : [0, 2, 3]
```

It can be noted that when a splitting is done, a new object is created, not a reference:

```
[90]: x = [1, 2, 3, 4]
      y = x[:2]
      y[0] = 0
      print("x : ", x)
      print("y : ", y)
```

```
x : [1, 2, 3, 4]
y : [0, 2]
```

4.1.2 Sorting

```
[1]: x = [2, 1, 4, 3]
      x.sort()
      print(x)
```

```
[1, 2, 3, 4]
```

4.2 Tuples

The tuples are sequences of Python objects. To create a tuple, one lists the values, separated by commas. Unlike lists, tuples are inalterable (i.e. cannot be modified after they have been created).

```
[92]: x = 1, 4, 9, 16, 25
      print(x)
      x[0] = 1
```

```
(1, 4, 9, 16, 25)
```

```
-----
TypeError                                Traceback (most recent call last)
Cell In[92], line 3
      1 x = 1, 4, 9, 16, 25
      2 print(x)
----> 3 x[0] = 1

TypeError: 'tuple' object does not support item assignment
```

4.3 Sets

Sets are unordered collections of unique elements. The sets are unalterable, not indexed. To create a set, Python provides the `set()` function. One or more elements constituting the set are provided, separated by commas and surrounded by braces. During the creation of a set, if there are duplicates in the values provided, these will be deleted to keep only one value

```
[93]: new_set = set({"Marseille", "Aix-en-Provence", "Nice", "Rennes"})  
      print(new_set)
```

```
{'Nice', 'Marseille', 'Aix-en-Provence', 'Rennes'}
```

Equivalently, rather than using the `set()` function, the set can only be defined using the brackets:

```
[ ]: new_set = {"Marseille", "Aix-en-Provence", "Nice", "Rennes"}  
     print(new_set)
```

```
[94]: new_set = set({"Marseille", "Aix-en-Provence", "Nice", "Marseille", "Rennes"})  
      print(new_set)
```

```
{'Nice', 'Marseille', 'Rennes', 'Aix-en-Provence'}
```

If the element is already present, it will not be added:

```
[95]: new_set.add("Rennes")  
      print(new_set)
```

```
{'Nice', 'Marseille', 'Rennes', 'Aix-en-Provence'}
```

```
[96]: new_set.add("Toulon")  
      print(new_set)
```

```
{'Nice', 'Toulon', 'Aix-en-Provence', 'Marseille', 'Rennes'}
```

```
[97]: new_set.remove("Toulon")  
      print(new_set)
```

```
{'Nice', 'Aix-en-Provence', 'Marseille', 'Rennes'}
```

```
[98]: print("Marseille" in new_set)
```

True

```
[99]: new_set = set({"Marseille", "Aix-en-Provence", "Nice"})  
      y = new_set.copy()  
      y.add("Toulon")  
      print("y : ", y)
```

```
y : {'Nice', 'Toulon', 'Marseille', 'Aix-en-Provence'}
```

Conversion to a List One of the interests of sets is that they contain unique elements. Also, when you want to obtain the distinct elements of a list, it is possible to convert it into a set (with the `set()` function), then to convert the set into a list (with the `list()` function):

```
[102]: my_list = ["Marseille", "Aix-en-Provence", "Marseille", "Marseille"]
       print(my_list)
```

```
['Marseille', 'Aix-en-Provence', 'Marseille', 'Marseille']
```

```
[103]: my_set = set(my_list)
       print(my_set)
```

```
{'Marseille', 'Aix-en-Provence'}
```

```
[104]: my_new_list = list(my_set)
       print(my_new_list)
```

```
['Marseille', 'Aix-en-Provence']
```

4.4 Dictionaries

Python dictionaries are an implementation of key-value objects, the keys being indexed. Keys are often text, values can be of different types and structures. To create a dictionary, you can proceed by using braces (`{}`).

```
[106]: my_dict = { "nom": "Kyrie",
                  "prenom": "John",
                  "naissance": 1992,
                  "equipes": ["Cleveland", "Boston"]}
       print(my_dict)
       print(my_dict["prenom"])
       print("age" in my_dict)
```

```
{'nom': 'Kyrie', 'prenom': 'John', 'naissance': 1992, 'equipes': ['Cleveland', 'Boston']}
```

```
John
```

```
False
```

```
[108]: the_keys = my_dict.keys()
       print(the_keys)
       the_keys_list = list(the_keys)
       print(the_keys_list)
```

```
dict_keys(['nom', 'prenom', 'naissance', 'equipes'])
['nom', 'prenom', 'naissance', 'equipes']
```

```
[109]: the_values = my_dict.values()
       print(the_values)
```



```
dict_values(['Kyrie', 'John', 1992, ['Cleveland', 'Boston']])
```

5 Operators

To raise a number to a given power, we use two stars (**) the inclusion tests are performed using the operator in.

```
[110]: print(3 in [1,2, 3])
       print(4 not in [1,2, 3])
       dictionnaire = {"nom": "Rockwell", "prenom": "Criquette"}
       "age" not in dictionnaire.keys()
```

True

True

```
[110]: True
```

And, Or

```
[111]: x = True
       y = False
       print(x and y)

       x = True
       y = True
       print(x and y)
```

False

True

Look list of math and stat functions

6 Loading and Saving Data

When we launch Jupyter Notebook, a tree structure is displayed, and we navigate inside it to create or open a notebook. The directory containing the notebook is the current directory. When Python is told to import data (or export objects), the origin (or destination) will be indicated relatively in the current directory, unless absolute paths (i.e., a path from the root /) are used.

If a Python program is started from a terminal, the current directory is the directory in which the terminal is located at the time the program is started.

```
[114]: import os
       cwd = os.getcwd()
       print(cwd)
       os.listdir()
```

```
/Users/yigitaydede/Library/CloudStorage/Dropbox/Python
```

```
[114]: ['PythR.Rmd',
        'try.py',
        '.Rhistory',
        'Untitled.ipynb',
        'PythR_v2.pdf',
        'PythR_v2.Rmd',
        'PythR.html',
        'StatisticsMachineLearningPython.pdf',
        '.ipynb_checkpoints',
        'flights.csv']
```

```
[118]: path = "flights.csv"
my_file = open(path, mode = "r")
print(my_file.read())
```

IOPub data rate exceeded.

The notebook server will temporarily stop sending output to the client in order to avoid crashing it.
To change this limit, set the config variable
`--NotebookApp.iopub_data_rate_limit`.

Current values:

NotebookApp.iopub_data_rate_limit=1000000.0 (bytes/sec)

NotebookApp.rate_limit_window=3.0 (secs)

It is important to remember to close the file once we have finished using it. To do this, we use the `close()` method.

```
[119]: my_file.close()
```

A common practice in Python is to open a file in a `with` block. The reason for this choice is that a file opened in such a block is automatically closed at the end of the block.

```
[120]: path = "flights.csv"
num_lines = 5 # Number of lines you want to print (similar to head in R)

with open(path, mode="r") as my_file:
    for i in range(num_lines):
        line = my_file.readline()
        print(line.strip()) # strip() removes trailing newlines or spaces
```

```
YEAR,MONTH,DAY,DAY_OF_WEEK,AIRLINE,FLIGHT_NUMBER,TAIL_NUMBER,ORIGIN_AIRPORT,DEST
INATION_AIRPORT,SCHEDULED_DEPARTURE,DEPARTURE_TIME,DEPARTURE_DELAY,TAXI_OUT,WHEE
LS_OFF,SCHEDULED_TIME,ELAPSED_TIME,AIR_TIME,DISTANCE,WHEELS_ON,TAXI_IN,SCHEDULED
_ARRIVAL,ARRIVAL_TIME,ARRIVAL_DELAY,DIVERTED,CANCELLED,CANCELLATION_REASON,AIR_S
YSTEM_DELAY,SECURITY_DELAY,AIRLINE_DELAY,LATE_AIRCRAFT_DELAY,WEATHER_DELAY
2015,1,1,4,AS,98,N407AS,ANC,SEA,0005,2354,-
11,21,0015,205,194,169,1448,0404,4,0430,0408,-22,0,0,,,,,
```

```

2015,1,1,4,AA,2336,N3KUA,LAX,PBI,0010,0002,-
8,12,0014,280,279,263,2330,0737,4,0750,0741,-9,0,0,,,,,
2015,1,1,4,US,840,N171US,SFO,CLT,0020,0018,-
2,16,0034,286,293,266,2296,0800,11,0806,0811,5,0,0,,,,,
2015,1,1,4,AA,258,N3HYAA,LAX,MIA,0020,0015,-
5,15,0030,285,281,258,2342,0748,8,0805,0756,-9,0,0,,,,,

```

In this code:

- `open(path, mode="r")` opens the file in read mode.
- `with` is used for context management. It handles opening and closing the file properly, even if an error occurs.
- `readline()` reads one line from the file each time it is called.
- The `for` loop runs `num_lines` times (5 in this case), each time printing a line from the file.
- `line.strip()` is used to remove any extra newline characters or spaces at the end of each line for cleaner output.

This code will print the first 5 lines of the file “flights.csv”. You can adjust `num_lines` to print as many lines as you need.

In Python, the most common way to import a CSV file as a `DataFrame` (which is similar to a `data.frame` in R) is to use the `pandas` library. `Pandas` is a powerful data manipulation and analysis tool, and it provides a `DataFrame` object that is quite similar to R’s `data.frame`.

Main Values for How to Open Files. `r`: Opening to read (default), `w`:Opening to write, `x`:Opening to create a document, fails if the file already exists, `a`: Opening to write, adding at the end of the file if it already exists, `+`: Opening for update (read and write) To be added to an opening mode for binary files (`rb` or `wb`), `t`:Text mode (automatic decoding of bytes in Unicode). Default if not specified (adds to the mode, like `b`)

```
[121]: import pandas as pd
```

```

path = "flights.csv"
df = pd.read_csv(path)

```

```
# Now you can work with the DataFrame 'df' just like a data.frame in R
```

```

/var/folders/wt/4xtk6v051vd349k3wktfld480000gn/T/ipykernel_14181/1726951512.py:4
: DtypeWarning: Columns (7,8) have mixed types. Specify dtype option on import
or set low_memory=False.
df = pd.read_csv(path)

```

The warning you’re encountering, `DtypeWarning: Columns (7,8) have mixed types. Specify dtype option on import or set low_memory=False`, is raised by `pandas` when it detects columns in your CSV file that have mixed data types. This can happen if, for example, a column contains both numbers and strings.

If you’re okay with `pandas` inferring the data types, and the mixed types don’t pose a problem for your analysis, you can choose to ignore this warning. After loading the data, if you find that certain columns need to be a different data type, you can convert them using the `astype()` method.

```
[122]: print(df.head())
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	\
0	2015	1	1	4	AS	98	N407AS	
1	2015	1	1	4	AA	2336	N3KUA	
2	2015	1	1	4	US	840	N171US	
3	2015	1	1	4	AA	258	N3HYAA	
4	2015	1	1	4	AS	135	N527AS	

	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	...	ARRIVAL_TIME	\
0	ANC	SEA	5	...	408.0	
1	LAX	PBI	10	...	741.0	
2	SFO	CLT	20	...	811.0	
3	LAX	MIA	20	...	756.0	
4	SEA	ANC	25	...	259.0	

	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON	AIR_SYSTEM_DELAY	\
0	-22.0	0	0	NaN	NaN	
1	-9.0	0	0	NaN	NaN	
2	5.0	0	0	NaN	NaN	
3	-9.0	0	0	NaN	NaN	
4	-21.0	0	0	NaN	NaN	

	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	WEATHER_DELAY
0	NaN	NaN	NaN	NaN
1	NaN	NaN	NaN	NaN
2	NaN	NaN	NaN	NaN
3	NaN	NaN	NaN	NaN
4	NaN	NaN	NaN	NaN

[5 rows x 31 columns]

When using `pandas.read_csv()` to read a CSV file in Python, you don't need to specify a file mode like 'r' for read or 'w' for write, as you would when using Python's built-in `open()` function. The `read_csv()` function from pandas is specifically designed for reading CSV files, and it handles the file opening and reading internally. It always opens the file in the appropriate mode for reading.

Import from the Internet To import a text file from the Internet, methods from the `urllib` library can be used:

```
[125]: import urllib
from urllib.request import urlopen
url = "http://egallic.fr/Enseignement/Python/fichiers_exemples/fichier_texte.
      ↪txt"
with urllib.request.urlopen(url) as my_file:
    data = my_file.read()
print(data)
```

```
b"Bonjour, je suis un fichier au format txt.\nJe contiens plusieurs lignes,  
l'id\x3\xa9e \xc3\xa9tant de montrer comment fonctionne l'importation d'un tel  
fichier dans Python.\nTrois lignes devraient suffir."
```

```
[126]: print(data.decode())
```

```
Bonjour, je suis un fichier au format txt.  
Je contiens plusieurs lignes, l'idée étant de montrer comment fonctionne  
l'importation d'un tel fichier dans Python.  
Trois lignes devraient suffir.
```

CSV Many databases export their data to CSV (e.g., World Bank, FAO, Eurostat, etc.). To import them into Python, you can use the `csv` module. Again, we use the `open()` function, with the parameters described earlier

```
[129]: import csv  
path = "flights.csv"  
with open(path) as my_file:  
    my_file_reader = csv.reader(my_file, delimiter=',', quotechar='"')  
    data = [x for x in my_file_reader]  
  
# Print the first five observations  
for row in data[:5]:  
    print(row)
```

```
['YEAR', 'MONTH', 'DAY', 'DAY_OF_WEEK', 'AIRLINE', 'FLIGHT_NUMBER',  
'TAIL_NUMBER', 'ORIGIN_AIRPORT', 'DESTINATION_AIRPORT', 'SCHEDULED_DEPARTURE',  
'DEPARTURE_TIME', 'DEPARTURE_DELAY', 'TAXI_OUT', 'WHEELS_OFF', 'SCHEDULED_TIME',  
'ELAPSED_TIME', 'AIR_TIME', 'DISTANCE', 'WHEELS_ON', 'TAXI_IN',  
'SCHEDULED_ARRIVAL', 'ARRIVAL_TIME', 'ARRIVAL_DELAY', 'DIVERTED', 'CANCELLED',  
'CANCELLATION_REASON', 'AIR_SYSTEM_DELAY', 'SECURITY_DELAY', 'AIRLINE_DELAY',  
'LATE_AIRCRAFT_DELAY', 'WEATHER_DELAY']  
['2015', '1', '1', '4', 'AS', '98', 'N407AS', 'ANC', 'SEA', '0005', '2354',  
'-11', '21', '0015', '205', '194', '169', '1448', '0404', '4', '0430', '0408',  
'-22', '0', '0', '', '', '', '', '', '']  
['2015', '1', '1', '4', 'AA', '2336', 'N3KUAA', 'LAX', 'PBI', '0010', '0002',  
'-8', '12', '0014', '280', '279', '263', '2330', '0737', '4', '0750', '0741',  
'-9', '0', '0', '', '', '', '', '', '']  
['2015', '1', '1', '4', 'US', '840', 'N171US', 'SFO', 'CLT', '0020', '0018',  
'-2', '16', '0034', '286', '293', '266', '2296', '0800', '11', '0806', '0811',  
'5', '0', '0', '', '', '', '', '', '']  
['2015', '1', '1', '4', 'AA', '258', 'N3HYAA', 'LAX', 'MIA', '0020', '0015',  
'-5', '15', '0030', '285', '281', '258', '2342', '0748', '8', '0805', '0756',  
'-9', '0', '0', '', '', '', '', '', '']
```

The two code snippets illustrate two different methods of reading a CSV file in Python: one using the built-in `csv` module (just above) and the other using the `pandas` library. Both methods achieve the same basic goal — reading data from a CSV file — but they have some key differences in

terms of ease of use, functionality, and the data structure they return. `pandas.read_csv` returns a `DataFrame`, a powerful data structure that provides a lot of functionalities for data manipulation and analysis.

Excel Files We can use the pandas library to read Excel files in Python. Pandas provides the `read_excel()` function, which is specifically designed for this purpose. This function can handle various Excel formats such as `.xlsx`, `.xslm`, `.xlsb`, and `.xls`. However, to read Excel files, we need to have the `openpyxl` (for `.xlsx` files) or `xlrd` (for older `.xls` files) packages installed. You can install these packages via `pip`:

`pip install openpyxl` and `pip install xlrd`

```
[136]: df2 = pd.read_excel("660218A.xlsx", header = 0)
print(df2.head())
```

		NAME	ID	50	/100	50.1	/100.1	40	\
0	Alimohammadi	Sagvand, Sanaz	A00422960	49.0	98	47.0	94	31.0	
1		Biao, Yanan	A00429012	48.5	97	48.0	96	38.0	
2		Cai, Ying	A00431323	49.5	99	49.5	99	28.5	
3		Cao, Hanmeng	A00406435	50.0	100	49.0	98	34.5	
4		Cao, Hanyu	A00398061	46.0	92	47.0	94	39.0	

	/100.2	60	/100.3	A25MT30F45	A25F75	Letter	Mark
0	77.50	46.0	76.666667	81.750	81.500		A-
1	95.00	57.0	95.000000	95.375	95.375		A+
2	71.25	46.0	76.666667	80.625	82.250		A-
3	86.25	59.5	99.166667	95.250	99.125		A+
4	97.50	58.0	96.666667	96.000	95.750		A+

Exporting In Python, especially in data science and analytics, the most common file types for saving work depend on the nature of the work and the requirements for data storage, sharing, and interoperability. Here are some of the most commonly used file types:

1. CSV (Comma-Separated Values): Widely used for storing tabular data. Easy to read and write, and can be opened by most spreadsheet programs like Microsoft Excel, Google Sheets, etc. Ideal for relatively simple datasets without complex structures.
2. Excel Files (`.xlsx` or `.xls`): Commonly used when sharing data with non-technical stakeholders or when additional formatting and data organization (like multiple sheets) are required.
3. Pickle Files (`.pkl`): Python-specific binary format. Used for saving Python objects, including `DataFrames`, preserving their data types and index structures. Not human-readable and not suitable for sharing data across different programming environments.
4. JSON (JavaScript Object Notation): Useful for storing data in a structured, hierarchical format. Human-readable and commonly used for web data and API interactions. Good for data that fits well into a nested, key-value structure.
5. Parquet and Feather Formats: Efficient, columnar storage formats. Ideal for large datasets, and used in big data processing and analytics. Feather is particularly useful for data interchange between Python and R.
6. HDF5 (Hierarchical Data Format version 5): Suitable for storing large quantities of scientific data. Supports data compression and can handle complex data structures.

7. **SQL Databases:** For projects involving relational databases, saving data directly to SQL databases (like SQLite, MySQL, PostgreSQL) is common. Useful for structured data that needs to be queried or joined with other data in a database.
8. **Python Scripts (.py) and Jupyter Notebooks (.ipynb):** For saving code, analysis, and documentation. Jupyter Notebooks are particularly popular for exploratory data analysis, as they allow you to combine code, text, and visualizations.

Saving your DataFrame `data` (or `df2`) as a CSV file using pandas is straightforward. You can use the `to_csv` method provided by pandas. Here's how you can do it with chunks (when you have large data, like `data`):

```
[143]: print(len(df))
```

5819079

```
[141]: chunk_size = 500000 # Size of each chunk - this is just an example
for i in range(0, len(data), chunk_size):
    df[i:i+chunk_size].to_csv(f'data_{i}.csv', index=False)
```

When working with data that's been split into multiple files or batches due to its size, and you need to perform search operations or apply conditions to select specific observations, you typically need to run the search or filter algorithm on each batch. Here's a general approach to how this can be done:

1. **Iterate Through Each Chunk:** Loop through each file (chunk) and perform the necessary search or filtering operation on each one. This can be done by loading each chunk into a pandas DataFrame, applying your conditions to filter or search the data, and then either processing the data immediately or storing the results for further use.
2. **Store or Aggregate Results:** As you process each chunk, you can either: Store the filtered/selected data from each chunk into a new file or a database for later use. Aggregate or summarize the data in memory, if feasible.
3. **Combine Results if Necessary:** After processing all chunks, you might end up with multiple smaller datasets (each representing the filtered data from one chunk). These can be combined, if necessary and if memory allows, or can be left as separate files for further processing.

```
[144]: import glob

# Condition/filter function
def filter_data(df):
    return df[df['DESTINATION_AIRPORT'] == "MIA"]

# List to hold the filtered data from each chunk
filtered_data = []

# Iterate through each chunk
for filename in glob.glob('data_*.csv'):
    chunk = pd.read_csv(filename)
    filtered_chunk = filter_data(chunk)
```

```

        filtered_data.append(filtered_chunk)

# Combine all filtered data into one DataFrame
final_data = pd.concat(filtered_data)

```

```

/var/folders/wt/4xtk6v051vd349k3wktfld480000gn/T/ipykernel_14181/51342753.py:12:
DtypeWarning: Columns (7,8) have mixed types. Specify dtype option on import or
set low_memory=False.

```

```

        chunk = pd.read_csv(filename)
/var/folders/wt/4xtk6v051vd349k3wktfld480000gn/T/ipykernel_14181/51342753.py:12:
DtypeWarning: Columns (7,8) have mixed types. Specify dtype option on import or
set low_memory=False.
        chunk = pd.read_csv(filename)

```

```
[145]: print(final_data.head())
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	\
305	2015	6	7	7	AA	1324	N3DRAA	
417	2015	6	7	7	AA	349	N016AA	
481	2015	6	7	7	UA	1615	N17229	
730	2015	6	7	7	AA	2471	N3GDAA	
834	2015	6	7	7	AA	919	N004AA	

	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	...	\
305	LGA	MIA	2100	...	
417	ATL	MIA	2110	...	
481	EWR	MIA	2115	...	
730	LAX	MIA	2145	...	
834	TPA	MIA	2155	...	

	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON	\
305	2357.0	-20.0	0	0	NaN	
417	2348.0	39.0	0	0	NaN	
481	13.0	-17.0	0	0	NaN	
730	604.0	12.0	0	0	NaN	
834	2238.0	-22.0	0	0	NaN	

	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	\
305	NaN	NaN	NaN	NaN	
417	0.0	0.0	39.0	0.0	
481	NaN	NaN	NaN	NaN	
730	NaN	NaN	NaN	NaN	
834	NaN	NaN	NaN	NaN	

	WEATHER_DELAY
305	NaN
417	0.0
481	NaN

730	NaN
834	NaN

[5 rows x 31 columns]

If you want to add an additional condition to your filtering process, you can modify the `filter_data` function to include this new condition. In addition to filtering rows where `DESTINATION_AIRPORT` is “MIA”, you also want to filter based on another condition, such as ‘`ORIGIN_AIRPORT`’ being “LAX”. You can use the logical AND operator (`&`) to combine these conditions in `pandas`.

```
[148]: # If you're confident that pandas' type inference won't adversely affect your
        ↪analysis,
        # you can choose to ignore this warning.

import warnings
warnings.filterwarnings('ignore', category=pd.errors.DtypeWarning)

def filter_data(df):
    condition1 = df['DESTINATION_AIRPORT'] == "MIA"
    condition2 = df['ORIGIN_AIRPORT'] == "LAX"
    return df[condition1 & condition2]

# List to hold the filtered data from each chunk
filtered_data = []

# Iterate through each chunk
for filename in glob.glob('data_*.csv'):
    chunk = pd.read_csv(filename)
    filtered_chunk = filter_data(chunk)
    filtered_data.append(filtered_chunk)

# Combine all filtered data into one DataFrame
final_data = pd.concat(filtered_data)

print(final_data.head())
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	\
730	2015	6	7	7	AA	2471	N3GDAA	
1166	2015	6	7	7	DL	1168	N3766	
1498	2015	6	7	7	AA	1538	N3KHAA	
1539	2015	6	8	1	AA	260	N3HMAA	
4126	2015	6	8	1	AA	68	N5EAAA	

	ORIGIN_AIRPORT	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	...	\
730	LAX	MIA	2145	...	
1166	LAX	MIA	2220	...	
1498	LAX	MIA	2355	...	
1539	LAX	MIA	10	...	

	LAX	MIA	735	...	
	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED	CANCELLATION_REASON \
730	604.0	12.0	0	0	NaN
1166	621.0	-3.0	0	0	NaN
1498	742.0	-19.0	0	0	NaN
1539	803.0	-19.0	0	0	NaN
4126	1519.0	-32.0	0	0	NaN
	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY	LATE_AIRCRAFT_DELAY	\
730	NaN	NaN	NaN	NaN	NaN
1166	NaN	NaN	NaN	NaN	NaN
1498	NaN	NaN	NaN	NaN	NaN
1539	NaN	NaN	NaN	NaN	NaN
4126	NaN	NaN	NaN	NaN	NaN
	WEATHER_DELAY				
730	NaN				
1166	NaN				
1498	NaN				
1539	NaN				
4126	NaN				

[5 rows x 31 columns]

6.1 Use R Magic in Python Notebooks:

If you primarily use Python but need to run some R code, you can use the `%R` magic command in a Python notebook. This requires the `rpy2` package in Python. Install `rpy2` via `pip install rpy2`. Then, at the start of your Python notebook, load the R magic:

```
[1]: %load_ext rpy2.ipython
```

```
[11]: %%R
library(tidyverse)

filter_data <- function(df) {
  df %>% filter(DESTINATION_AIRPORT == "MIA", ORIGIN_AIRPORT == "LAX")
}

read_csv_quietly <- function(filename) {
  read_csv(filename, show_col_types = FALSE)
}

# Get the list of files
# list.files() function with the apptern argument is to list files in the CD
↳ that
# matches a specific pattern
```

```

files <- list.files(pattern = "data_.*\\.csv")

# Read, filter, and combine data from each file
final_data <- files %>%
  map(read_csvquietly) %>%
  map(filter_data) %>%
  bind_rows()

# Display the head of the final data frame
head(final_data)

# A tibble: 6 × 31
  YEAR MONTH DAY DAY_OF_WEEK AIRLINE FLIGHT_NUMBER TAIL_NUMBER ORIGIN_AIRPORT
  <dbl> <dbl> <dbl>         <dbl> <chr>          <dbl> <chr>          <chr>
1  2015     1     1             4 AA              258 N3HYAA        LAX
2  2015     1     1             4 AA              115 N3CTAA        LAX
3  2015     1     1             4 AA              306 N859AA        LAX
4  2015     1     1             4 AA              208 N355AA        LAX
5  2015     1     1             4 AA              124 N7BEAA        LAX
6  2015     1     1             4 AA               28 N358AA        LAX
# 23 more variables: DESTINATION_AIRPORT <chr>, SCHEDULED_DEPARTURE <dbl>,
# DEPARTURE_TIME <dbl>, DEPARTURE_DELAY <dbl>, TAXI_OUT <dbl>,
# WHEELS_OFF <dbl>, SCHEDULED_TIME <dbl>, ELAPSED_TIME <dbl>, AIR_TIME <dbl>,
# DISTANCE <dbl>, WHEELS_ON <dbl>, TAXI_IN <dbl>, SCHEDULED_ARRIVAL <dbl>,
# ARRIVAL_TIME <dbl>, ARRIVAL_DELAY <dbl>, DIVERTED <dbl>, CANCELLED <dbl>,
# CANCELLATION_REASON <chr>, AIR_SYSTEM_DELAY <dbl>, SECURITY_DELAY <dbl>,
# AIRLINE_DELAY <dbl>, LATE_AIRCRAFT_DELAY <dbl>, WEATHER_DELAY <dbl>

```

In the code snippet provided above, `map` is a function from the `purrr` package, which is part of the tidyverse suite of packages in R. The `map` function is used to apply a function to each element of a list or vector and is particularly useful for performing operations on lists in a concise and readable way.

```

[18]: %>%R

# OR

library(tidyverse)

filter_data <- function(dff){
  dff %>% filter(DESTINATION_AIRPORT == "MIA", ORIGIN_AIRPORT == "LAX")
}

# matches a specific pattern
files <- list.files(pattern = "data_.*\\.csv")

filtered_data <- NULL

```

```
# Iterate through each chunk
for (filename in files) {
  chunk <- read_csv(filename, show_col_types = FALSE)
  filtered_chunk <- filter_data(chunk)
  filtered_data <- rbind(filtered_data, filtered_chunk)
}

head(as.data.frame(filtered_data))
```

	YEAR	MONTH	DAY	DAY_OF_WEEK	AIRLINE	FLIGHT_NUMBER	TAIL_NUMBER	ORIGIN_AIRPORT						
1	2015	1	1	4	AA	258	N3HYAA	LAX						
2	2015	1	1	4	AA	115	N3CTAA	LAX						
3	2015	1	1	4	AA	306	N859AA	LAX						
4	2015	1	1	4	AA	208	N355AA	LAX						
5	2015	1	1	4	AA	124	N7BEAA	LAX						
6	2015	1	1	4	AA	28	N358AA	LAX						
	DESTINATION_AIRPORT	SCHEDULED_DEPARTURE	DEPARTURE_TIME	DEPARTURE_DELAY										
1	MIA	20	15	-5										
2	MIA	105	103	-2										
3	MIA	800	759	-1										
4	MIA	840	842	2										
5	MIA	1250	1305	15										
6	MIA	1500	1456	-4										
	TAXI_OUT	WHEELS_OFF	SCHEDULED_TIME	ELAPSED_TIME	AIR_TIME	DISTANCE	WHEELS_ON							
1	15	30	285	281	258	2342	748							
2	14	117	286	276	255	2342	832							
3	13	812	292	280	257	2342	1529							
4	19	901	293	285	264	2342	1625							
5	17	1322	285	277	256	2342	2038							
6	19	1515	290	305	284	2342	2259							
	TAXI_IN	SCHEDULED_ARRIVAL	ARRIVAL_TIME	ARRIVAL_DELAY	DIVERTED	CANCELLED								
1	8	805	756	-9	0	0								
2	7	851	839	-12	0	0								
3	10	1552	1539	-13	0	0								
4	2	1633	1627	-6	0	0								
5	4	2035	2042	7	0	0								
6	2	2250	2301	11	0	0								
	CANCELLATION_REASON	AIR_SYSTEM_DELAY	SECURITY_DELAY	AIRLINE_DELAY										
1	<NA>	NA	NA	NA										
2	<NA>	NA	NA	NA										
3	<NA>	NA	NA	NA										
4	<NA>	NA	NA	NA										
5	<NA>	NA	NA	NA										
6	<NA>	NA	NA	NA										
	LATE_AIRCRAFT_DELAY	WEATHER_DELAY												
1	NA	NA												
2	NA	NA												

3	NA	NA
4	NA	NA
5	NA	NA
6	NA	NA

7 Conditions

Often, depending on the evaluation of an expression, one wants to perform one operation rather than another.

7.1 if

if expression: instruction

```
[19]: x = 2
      if x == 2:
          print("Hello")
```

Hello

```
[20]: # Not this
      x=3
      if x == 2:
          print("Hello")
```

Inside the block, several instructions can be written that will be evaluated if the expression is True:

```
[21]: x=2
      if x == 2:
          y = "Hello"
          print(y + ", x is : " + str(x))
```

Hello, x is : 2

7.2 if-else

if expression: instructions else: other_instruction

```
[22]: temperature = 26
      heat = ""

      if temperature > 28:
          heat = "hot"
      else:
          heat = "cold"

      print("It is " + heat + " out there")
```

It is cold out there

7.3 if-elif

```
[23]: temperature = -4
      heat = ""

      if temperature > 28:
          heat = "hot"
      elif temperature <= 28 and temperature > 15:
          heat = "not too hot"
      elif temperature <= 15 and temperature > 0:
          heat = "cold"
      else:
          heat = "very cold"

      print("It is " + heat + " out there")
```

It is very cold out there

8 Loops

When the same operation has to be repeated several times, for a given number of times or as long as a condition is verified (or as long as it is not verified), loops can be used, which is much less painful than evaluating by hand or by copying and pasting the same instruction.

We will discuss two types of loops in this chapter:

- those for which we do not know a priori the number of iterations (the number of repetitions) to be performed: `while()` loops
- those for which we know a priori how many iterations are necessary: `for()` loops

```
[24]: x = 100
      while x/3 > 1:
          print(x/3)
          x = x/3
```

33.333333333333336
11.111111111111112
3.703703703703704
1.234567901234568

```
[28]: # Single loop
      total = 0
      for value in [3, 2, 19]:
          total += value

      print('Total is: {}'.format(total))
```

Total is: 24

We also took advantage of the increment notation in Python: the expression `a += b` increment is equivalent to `a = a + b`. R doesn't have an in-built shorthand operator for incrementing a variable. So, the expression `a += b` that you would use in other languages is written as `a = a + b` in R.

```
[29]: # See the difference

total = 0
for value in [3, 2, 19]:
    total += value
    print('Total is: {}'.format(total))
```

```
Total is: 3
Total is: 5
Total is: 24
```

```
[30]: ## Double loop
total = 0
for value in [2,3,19]:
    for weight in [3, 2, 1]:
        total += value * weight

print('Total is: {}'.format(total))
```

```
Total is: 144
```

When we know the number of iterations in advance, we can use a `for()` loop. The syntax is as follows:

```
[25]: message = "The squared value of {} is {}"
n = 10
for i in range(0, n + 1):
    print(message.format(i, i ** 2))
```

```
The squared value of 0 is 0
The squared value of 1 is 1
The squared value of 2 is 4
The squared value of 3 is 9
The squared value of 4 is 16
The squared value of 5 is 25
The squared value of 6 is 36
The squared value of 7 is 49
The squared value of 8 is 64
The squared value of 9 is 81
The squared value of 10 is 100
```

If we want to store the result in a list:

```
[26]: n=10
n_squares = []

for i in range(0, n+1):
    n_squares.append(i**2)

print(n_squares)
```

```
[0, 1, 4, 9, 16, 25, 36, 49, 64, 81, 100]
```

```
[27]: %%R

n <- 10
n_squares <- c() # Initialize an empty vector

# R starts at 1 not zero:-)
for (i in 0:n) {
    n_squares <- c(n_squares, i^2) # Append the square of i to the vector
}

print(n_squares)
```

```
[1] 0 1 4 9 16 25 36 49 64 81 100
```

```
[30]: message = "There is(are) {} letter(s) in the name: {}"
for first_name in ["Pascaline", "Gauthier", "Xuan", "Jimmy"]:
    print(message.format(len(first_name), first_name))

# Cleaner

message = "There is(are) {} letter(s) in the name: {}"
for i in ["Pascaline", "Gauthier", "Xuan", "Jimmy"]:
    print(message.format(len(i), i))
```

```
There is(are) 9 letter(s) in the name: Pascaline
There is(are) 8 letter(s) in the name: Gauthier
There is(are) 4 letter(s) in the name: Xuan
There is(are) 5 letter(s) in the name: Jimmy
There is(are) 9 letter(s) in the name: Pascaline
There is(are) 8 letter(s) in the name: Gauthier
There is(are) 4 letter(s) in the name: Xuan
There is(are) 5 letter(s) in the name: Jimmy
```

Here is the R code, where `sprintf()` is used to format the message. It replaces `%s` with the number of characters in the name (`nchar(first_name)`) and the name itself (`first_name`).

```
[32]: %%R
```



```

message <- "There is(are) %s letter(s) in the name: %s"
first_names <- c("Pascaline", "Gauthier", "Xuan", "Jimmy")

for (i in first_names) {
  formatted_message <- sprintf(message, nchar(i), i)
  print(formatted_message)
}

# OR

message <- "There is(are) %s letter(s) in the name: %s"

for (i in c("Pascaline", "Gauthier", "Xuan", "Jimmy")) {
  formatted_message <- sprintf(message, nchar(i), i)
  print(formatted_message)
}

```

```

[1] "There is(are) 9 letter(s) in the name: Pascaline"
[1] "There is(are) 8 letter(s) in the name: Gauthier"
[1] "There is(are) 4 letter(s) in the name: Xuan"
[1] "There is(are) 5 letter(s) in the name: Jimmy"
[1] "There is(are) 9 letter(s) in the name: Pascaline"
[1] "There is(are) 8 letter(s) in the name: Gauthier"
[1] "There is(are) 4 letter(s) in the name: Xuan"
[1] "There is(are) 5 letter(s) in the name: Jimmy"

```

```

[35]: # This will be interesting for more complex loops
import time # Importing time for demonstration purposes

for i in range(3):
    for j in range(3):
        print(f"i equals {i} and j equals {j}", end="\r")
        time.sleep(1) # Adding a delay for demonstration purposes

```

i equals 2 and j equals 2

```

[39]: %%R

# Try it in RStudio! Doesn't work here:-)

library(progress)

# Define the total number of iterations
total <- 3 * 3 # As we have two loops running from 0 to 2

# Create a progress bar
pb <- progress_bar$new(

```

```

    format = "[:bar] :percent :elapsedfull",
    total = total,
    clear = FALSE,
    width = 60
)

# Iterate with a progress bar
for (i in 0:2) {
  for (j in 0:2) {
    # Update the progress bar
    pb$tick()

    # Your loop content goes here
    # ...

    Sys.sleep(0.1) # Added only to slow down the loop for demonstration
  }
}

```

```

[40]: import ipywidgets as widgets
from IPython.display import display
import time

progress = widgets.IntProgress(value=0, min=0, max=9, description='Loading:')
display(progress)

for i in range(9):
    time.sleep(0.5) # Simulating work
    progress.value += 1

```

IntProgress(value=0, description='Loading:', max=9)

We don't need to memorize these, but when we need them, we know they exist!

```

[41]: message = "New value for j: {}"
j = 10
for i in range(0, 4):
    j += 5
    print(message.format(j))

```

```

New value for j: 15
New value for j: 20
New value for j: 25
New value for j: 30

```

In a loop, if we want to increment a counter, we can use the symbol += rather than writing 'counter = counter + . . . '. But this is not available in R

```
[42]: %%R

j <- 10
for (i in 1:4) {
  j <- j + 5
  cat("New value for j:", j, "\n")
}
```

```
New value for j: 15
New value for j: 20
New value for j: 25
New value for j: 30
```

Here is a good example: Choose a ‘mystery’ number between 1 and 100, and store it in an object called `mystere_number`. Then, create a loop that at each iteration performs a random draw of an integer between 1 and 100. As long as the number drawn is different from the mystery number, the loop must continue. At the output of the loop, a variable called `nb_drawings` will contain the number of prints made to obtain the mystery number.

```
[64]: %%R

mystere_number <- sample(1:100, 1)
rd <- 0

while (rd != mystere_number) {
  rd <- sample(1:100, 1)
}

c(mystere_number, nb_drawings)
```

```
[1] 39 39
```

```
[94]: import random
import ipywidgets as widgets
from IPython.display import display
import time

progress = widgets.IntProgress(value=0, min=0, max=3000, description='Loading:')
display(progress)

mystere_number = random.randint(0, 2000)
rd = 0

while rd != mystere_number:
  rd = random.randint(0, 2000)
  progress.value += 1
```

```
print(mystere_number, rd)
```

```
IntProgress(value=0, description='Loading:', max=3000)
```

```
1746 1746
```

Use a loop to scan integers from 1 to 20 using a for loop, displaying in the console at each iteration if the current number is even.

```
[96]: message1 = "The value {} is even"
      message2 = "The value {} is not even"

      for i in range(1, 21):
          if i % 2 == 0:
              print(message1.format(i))
          else:
              print(message2.format(i))
```

```
The value 1 is not even
The value 2 is even
The value 3 is not even
The value 4 is even
The value 5 is not even
The value 6 is even
The value 7 is not even
The value 8 is even
The value 9 is not even
The value 10 is even
The value 11 is not even
The value 12 is even
The value 13 is not even
The value 14 is even
The value 15 is not even
The value 16 is even
The value 17 is not even
The value 18 is even
The value 19 is not even
The value 20 is even
```

```
[97]: %%R

message1 <- "The value %s is even"
message2 <- "The value %s is not even"

for (i in 1:20) {
  if (i %% 2 == 0) {
    cat(sprintf(message1, i), "\n")
  } else {
    cat(sprintf(message2, i), "\n")
  }
}
```

```

    }
}

# OR

for (i in 1:20) {
  if (i %% 2 == 0) {
    paste("The value ", i, " is even")
  } else {
    paste("The value ", i, " is not even")
  }
}

```

```

The value 1 is not even
The value 2 is even
The value 3 is not even
The value 4 is even
The value 5 is not even
The value 6 is even
The value 7 is not even
The value 8 is even
The value 9 is not even
The value 10 is even
The value 11 is not even
The value 12 is even
The value 13 is not even
The value 14 is even
The value 15 is not even
The value 16 is even
The value 17 is not even
The value 18 is even
The value 19 is not even
The value 20 is even

```

Perhaps a more common task would be to sum over (value, weight) pairs. For instance, to compute the average value of a random variable that takes on possible values 2, 3 or 19 with probability 0.2, 0.3, 0.5 respectively we would compute the weighted sum. Tasks such as this can often be accomplished using the `zip()` function that loops over a sequence of tuples.

```

[31]: total = 0
      for value, weight in zip([2,3,19],
                               [0.2,0.3,0.5]):
        total += weight * value
      print('Weighted average is: {0}'.format(total))

```

```

Weighted average is: 10.8

```

```
[33]: %%R

values <- c(2, 3, 19)
weights <- c(0.2, 0.3, 0.5)

total <- 0
for (i in 1:length(values)) {
  total <- total + weights[i] * values[i]
}

cat('Weighted average is:', total, '\n')
```

Weighted average is: 10.8

Wow, how stupid it's! A Simple version would be you'd perform a vectorized operation in NumPy

```
[32]: import numpy as np

values = np.array([2, 3, 19])
weights = np.array([0.2, 0.3, 0.5])

total = np.sum(weights * values)
print('Weighted average is:', total)
```

Weighted average is: 10.8

```
[34]: %%R

values <- c(2, 3, 19)
weights <- c(0.2, 0.3, 0.5)

total <- sum(weights * values)

cat('Weighted average is:', total, '\n')
```

Weighted average is: 10.8

9 FUNCTIONS

Most of the time, we use the basic functions or those contained in modules. However, when retrieving data online or formatting data imported from various sources, it may be necessary to create our own functions.

```
def name_function(arguments):      body of the function4
```

Once the function is defined, it is called by referring to its name: `name_function()`

```
[98]: def square(x):
      return x**2
```

```
square(2)
```

[98]: 4

```
[99]: %%R  
  
squarer <- function(x){  
  return(x^2)  
}  
squarer(2)
```

[1] 4

In both Python and R, functions can have positional arguments and keyword (named) arguments. However, there are some differences in how these arguments are handled in each language.

Python: **Positional Arguments:** These are arguments that can be called by their position in the function definition. **Keyword Arguments:** Also known as named arguments, these are specified by name. One key feature in Python is that you can have default values for keyword arguments. Example in Python:

```
[100]: def greet(name, msg="Hello"):  
           print(msg, name)  
  
greet("Alice")           # Uses default value for msg  
greet("Bob", "Goodbye")  # Overrides default value for msg  
greet(msg="Hi", name="Eve") # Using named arguments
```

Hello Alice
Goodbye Bob
Hi Eve

```
[101]: %%R  
  
greet <- function(name, msg="Hello") {  
  cat(msg, name, "\n")  
}  
  
greet("Alice")           # Uses default msg  
greet("Bob", "Goodbye")  # Overrides default msg  
greet(name="Eve", msg="Hi") # Named arguments, order can be changed  
greet("Charlie", msg="Hey") # Mix of positional and named arguments
```

Hello Alice
Goodbye Bob
Hi Eve
Hey Charlie

A function can be provided as an argument to another function.

```
[103]: def square(x):  
        """Returns the squared value of x"""  
        return x**2  
  
        def apply_fun_to_4(fun):  
            """Applies the function `fun` to 4"""  
            return fun(4)  
  
        print(apply_fun_to_4(square))
```

16

```
[110]: %%R  
  
# Define the square function  
square <- function(x) {  
  return(x^2)  
}  
  
# Define a function that applies another function to 4  
apply_fun_to_4 <- function(fun) {  
  return(fun(4))  
}  
  
# Apply the square function to 4 and print the result  
result <- apply_fun_to_4(square)  
cat("Result:", result, "\n")
```

Result: 16

[]: When a function **is** called, the body of that function **is** interpreted. Variables
↳ that have been defined **in** the body of the function are assigned to a local
↳ namespace.
In other words, they live only within this local space, which **is** created at the
↳ moment of the call of the function **and** destroyed at the end of it.

```
[106]: value = 10  
  
def f(x):  
    value = 2  
    return x + value  
  
f(5)
```

[106]: 7

If a variable is not defined in the body of a function, Python will search in a parent environment.

```
[108]: def f(y):  
        return y + value  
  
f(5)
```

[108]: 15

```
[116]: %%R  
  
# Same for R  
  
value = 10  
  
f <- function(x) {  
  value = 2  
  return(x + value)  
}  
  
# But  
  
k <- function(y) {  
  return(y + value)  
}  
  
c(f(5), value, k(5))
```

[1] 7 10 15

Python offers what are called lambdas functions, or anonymous functions. A lambda function has only one instruction whose result is that of the function.

```
[117]: def square(x):  
        return x**2  
  
## The same function with lamda  
square_2 = lambda x: x**2  
print(square_2(4))
```

16

It can sometimes be convenient to return several elements in return for a function. Although the list is a candidate for this feature, it may be better to use a dictionary, to be able to access the values with their key!

```
[122]: import statistics
```

```
def stat_des(x):
    """Returns the mean and standard deviation of `x`"""
    return {"mean": statistics.mean(x), "std_dev": statistics.stdev(x)}

x = [1,3,2,6,4,1,8,9,3,2]
res = stat_des(x)
print(res)

print(res["mean"])
```

```
{'mean': 3.9, 'std_dev': 2.8460498941515415}
3.9
```

```
[123]: %%R

# Define the function
stat_des <- function(x) {
  return(list(mean = mean(x), std_dev = sd(x)))
}

# Example data
x <- c(1, 3, 2, 6, 4, 1, 8, 9, 3, 2)

# Get the result as a named list
res <- stat_des(x)
print(res)

# Access elements of the list
print(res$mean) # Access the mean
print(res$std_dev) # Access the standard deviation
```

```
$mean
[1] 3.9

$std_dev
[1] 2.84605

[1] 3.9
[1] 2.84605
```

10 Introduction to Numpy

An important library for numerical calculations: NumPy (abbreviation of Numerical Python). It is common practice to import NumPy by assigning it the alias np: `import numpy as np`

10.1 Arrays

NumPy offers a popular data structure, arrays, on which calculations can be performed efficiently. In `numpy`, an array is a generic term for a multidimensional set of numbers. We use the `np.array()` function to define `x` and `y`, which are one-dimensional arrays, i.e. vectors.

```
[2]: import numpy as np

x = np.array([3, 4, 5])
y = np.array([4, 9, 7])

x + y
```

```
[2]: array([ 7, 13, 12])
```

In `numpy`, matrices are typically represented as two-dimensional arrays, and vectors as one-dimensional arrays. We can create a two-dimensional array as follows.

```
[3]: x = np.array([[1, 2], [3, 4]])
x
```

```
[3]: array([[1, 2],
           [3, 4]])
```

The object `x` has several attributes, or associated objects. To access an attribute of `x`, we type `x.attribute`, where we replace `attribute` with the name of the attribute. For instance, we can access the `ndim` attribute of `x` as follows.

```
[6]: print(x.ndim)
print(x.dtype)
```

```
2
int64
```

Why is `x` comprised of integers? This is because we created `x` by passing in exclusively integers to the `np.array()` function. If we had passed in any decimals, then we would have obtained an array of floating point numbers (i.e. real-valued numbers).

```
[7]: np.array([[1, 2], [3.0, 4]]).dtype
```

```
[7]: dtype('float64')
```

The array `x` is two-dimensional. We can find out the number of rows and columns by looking at its `shape` attribute

```
[8]: x.shape
```

```
[8]: (2, 2)
```

A method is a function that is associated with an object. For instance, given an array `x`, the expression `x.sum()` sums all of its elements, using the `sum()` method for arrays. The call `x.sum()` automatically provides `x` as the first argument to its `sum()` method.

```
[9]: x = np.array([1, 2, 3, 4])
     x.sum()
```

```
[9]: 10
```

We could also sum the elements of `x` by passing in `x` as an argument to the `np.sum()` function.

```
[10]: x = np.array([1, 2, 3, 4])
      np.sum(x)
```

```
[10]: 10
```

Pay attention here:

```
[13]: x = np.array([1, 2, 3, 4, 5, 6])
      print(x)

      x_reshaped = x.reshape((2, 3))
      print(x_reshaped)

      #Change the first element to 5
      x_reshaped[0, 0] = 5

      print(x_reshaped)
      print(x)
```

```
[1 2 3 4 5 6]
[[1 2 3]
 [4 5 6]]
[5 2 3 4 5 6]
[[5 2 3]
 [4 5 6]]
```

To our surprise, we discover that the first element of `x` has been modified as well! Modifying `x_reshaped` also modified `x` because the two objects occupy the same space in memory.

10.2 Random data

In data analytics, we often want to generate random data. the `np.random.normal()` function generates a vector of random normal variables. We can learn more about this function by looking at the help page, via a call to `np.random.normal?`. The first line of the help page reads `normal(loc=0.0, scale=1.0, size=None)`. This signature line tells us that the function's arguments are `loc`, `scale`, and `size`. These are keyword arguments, which means that when they are passed into the function, they can be referred to by name (in any order). By default, this function will generate random

normal variable(s) with mean (`loc`) 0 and standard deviation (`scale`) 1; further more, a single random variable will be generated unless the argument to size is changed.

```
[14]: x = np.random.normal(size=50)
      x
```

```
[14]: array([-3.98758621e-01, -1.32585520e+00,  2.99086619e-01, -9.00467046e-01,
        -6.24377293e-01,  1.38433653e+00, -1.98727515e+00, -1.33082457e-01,
        -9.71992791e-01,  6.63398551e-01,  3.46517153e-01, -1.63397597e+00,
         3.33672263e-01, -9.34686339e-01, -1.41242974e+00, -1.34432196e-01,
        -1.72489169e+00,  1.60784233e-01, -1.05829487e+00, -2.18321556e-02,
         1.89588479e-01, -1.39836483e-01, -8.93383689e-01, -7.83250154e-02,
         1.17285345e+00, -3.39760151e-01,  7.51434322e-02,  1.19200948e+00,
        -5.65394550e-01, -7.84100601e-01, -9.63238726e-02,  1.05682617e+00,
        -2.19975770e-03, -3.12787509e-01, -4.09556939e-01, -4.53440641e-03,
        -5.26438184e-01, -3.48628682e+00,  5.96224387e-01,  1.29367128e+00,
        -3.66914424e-01,  1.05576017e+00,  1.34325270e+00,  1.01442394e+00,
         1.16487534e+00, -2.99751323e-01,  1.61316775e+00,  3.66836276e-01,
        -2.25382573e-01,  1.15627344e-02])
```

We create an array `y` by adding an independent $N(50,1)$ random variable to each element of `x`.

```
[17]: y = x + np.random.normal(loc=50, scale=1, size=50)
      y
```

```
[17]: array([50.91313043, 48.35385782, 50.80903044, 50.54994561, 48.69165226,
        50.60932749, 48.78877271, 50.15383787, 48.42220353, 49.80142401,
        51.68320569, 47.47239158, 50.91925555, 49.34482161, 49.16704937,
        49.48369408, 47.72670698, 50.40858947, 49.02627041, 51.47861593,
        48.97465266, 49.92856398, 50.42015717, 50.57461143, 52.15778099,
        47.33148267, 49.79171346, 51.46796316, 50.55562685, 47.51896059,
        50.43897447, 51.001371  , 51.77533076, 49.44005047, 50.68029124,
        51.48579121, 50.05814013, 47.42541796, 50.95767462, 51.35853017,
        48.78070351, 49.77239821, 50.46416671, 50.91403132, 53.04549545,
        50.24357539, 51.88093265, 50.81059912, 49.5126853 , 51.05077079])
```

In order to ensure that our code provides exactly the same results each time it is run, we can set a random seed using the `np.random.default_rng()`

```
[24]: rng = np.random.default_rng(1303)
      print(rng.normal(scale=5, size=2))

      rng2 = np.random.default_rng(1303)
      print(rng2.normal(scale=5, size=2))
```

```
[ 4.09482632 -1.07485605]
[ 4.09482632 -1.07485605]
```

```
[27]: # Some statistical operations
rng = np.random.default_rng(3)
y = rng.standard_normal(10)
np.mean(y), y.mean()
```

```
[27]: (-0.1126795190952861, -0.1126795190952861)
```

```
[28]: #And
np.var(y), y.var(), np.mean((y - y.mean())**2)
```

```
[28]: (2.7243406406465125, 2.7243406406465125, 2.7243406406465125)
```

10.3 Column/Row operations

The `np.mean()`, `np.var()`, and `np.std()` functions can also be applied to the rows and columns of a matrix. To see this, we construct a 10×3 matrix of $N(0,1)$ random variables, and consider computing its row sums.

```
[29]: X = rng.standard_normal((10, 3))
X
```

```
[29]: array([[ 0.22578661, -0.35263079, -0.28128742],
        [-0.66804635, -1.05515055, -0.39080098],
        [ 0.48194539, -0.23855361,  0.9577587 ],
        [-0.19980213,  0.02425957,  1.54582085],
        [ 0.54510552, -0.50522874, -0.18283897],
        [ 0.54052513,  1.93508803, -0.26962033],
        [-0.24355868,  1.0023136 , -0.88645994],
        [-0.29172023,  0.88253897,  0.58035002],
        [ 0.0915167 ,  0.67010435, -2.82816231],
        [ 1.02130682, -0.95964476, -1.66861984]])
```

```
[30]: X.mean(axis=0)
# Sum of its rows by columns
```

```
[30]: array([ 0.15030588,  0.14030961, -0.34238602])
```

```
[31]: X.mean(axis=1)
# Sum of its columns by rows
```

```
[31]: array([-0.13604387, -0.70466596,  0.40038349,  0.45675943, -0.04765406,
          0.73533095, -0.04256834,  0.39038958, -0.68884708, -0.53565259])
```

10.4 Speed

Arrays are a useful structure for performing basic statistical operations as well as pseudo-random number generation. The structure of the tables is similar to that of the lists, but the latter are slower to process and use more memory. To be convinced:

```
[19]: # first lists

from random import random
from operator import truediv
l1 = [random() for i in range(1000)]
l2 = [random() for i in range(1000)]
%timeit s = sum(map(truediv,l1,l2))
```

24.6 μ s \pm 328 ns per loop (mean \pm std. dev. of 7 runs, 10,000 loops each)

To write the Python expression `l1 = [random() for i in range(1000)]` using a traditional for loop, you would first initialize an empty list and then append a random number to this list in each iteration of the loop. Here's how you can do it:

```
[23]: from random import random

l1 = [] # Initialize an empty list
for i in range(1000):
    l1.append(random()) # Append a random number to the list in each iteration
```

The output tells you that the code was very quick to execute, with an average time of 26.8 microseconds per loop, and the execution time was quite consistent, with a standard deviation of about 828 nanoseconds. The measurement was based on a total of 70,000 executions of the loop (10,000 loops per run over 7 runs).

```
[6]: %%R

library(microbenchmark)

# Generate random numbers
l1 <- runif(1000)
l2 <- runif(1000)

# Time the operation
timing_result <- microbenchmark(
  s <- sum(l1 / l2),
  times = 100L # Number of times to run the expression
)

print(timing_result)
```

Unit: microseconds

	expr	min	lq	mean	median	uq	max	neval
s <- sum(l1/l2)		2.764	2.8045	3.20613	2.8495	2.9435	17.567	100

The R results show that the operation is very quick, with most executions taking between 3.4 to 5.8 microseconds, but with some variability as indicated by the range from the minimum to the maximum times. The mean and median are close in value, suggesting a relatively symmetric distribution of execution times.

`l1 <- runif(1000)` in R is more concise compared to the equivalent Python code `l1 = [random() for i in range(1000)]`. This difference in conciseness can often be seen between R and Python, especially in tasks related to data generation and statistical operations.

R is a language designed specifically for statistical computing and has many built-in functions that are highly optimized for such tasks. Functions like `runif` provide a very straightforward and concise way to generate random numbers, which is a common requirement in statistical analysis and simulation.

Python, while extremely versatile and powerful, is a general-purpose programming language. Operations like generating random numbers require either using list comprehensions (as in your example) or loops, along with the standard library functions like `random.random()`. Python's approach is more verbose but also very flexible, allowing for complex operations within list comprehensions and loops.

Each language has its strengths, and in this case, R's specialized nature for statistical tasks makes it more concise for generating a sequence of random numbers.

Now, let's transform the two lists into NumPy tables with the `array()` method, and do the same calculation with a NumPy method:

```
[8]: a1 = np.array(l1)
      a2 = np.array(l2)
      %timeit s = np.sum(a1/a2)
```

5.6 μ s \pm 36 ns per loop (mean \pm std. dev. of 7 runs, 100,000 loops each)

As can be seen by executing these codes in an IPython environment, the execution time is much faster with the NumPy methods for this calculation. The creation of an array can be done with the `array()` method, from a list, as we just did:

```
[13]: list = [1,2,4]
      table = np.array(list)
      print(table)

      # OR

      table2 = np.array([1,2,4])
      print(table2)

      # Or

      list_2 = [ [1,2,3], [4,5,6] ]
      table_2 = np.array(list_2)
      print(table_2)

      print(type(table_2))
```

```
[1 2 4]
[1 2 4]
[[1 2 3]
```



```
[4 5 6]]
<class 'numpy.ndarray'>
```

Arrays in NumPy (Numerical Python) are commonly used for creating and working with matrices in Python. NumPy provides a powerful array object, which is a multi-dimensional (n-dimensional) homogeneous array. It's a central data structure of the NumPy library.

NumPy arrays can be used to represent both vectors (1D arrays) and matrices (2D arrays), as well as higher-dimensional data structures. For matrix operations, 2D arrays are typically used. Here's a basic example of how you can create and manipulate matrices with NumPy:

```
[14]: # Create a 2x3 matrix
matrix = np.array([[1, 2, 3], [4, 5, 6]])
print(matrix)

# Perform matrix operations
transpose = matrix.T # Transpose of the matrix
print("Transpose:\n", transpose)

product = np.dot(matrix, transpose) # Matrix multiplication
print("Matrix Product:\n", product)
```

```
[[1 2 3]
 [4 5 6]]
Transpose:
[[1 4]
 [2 5]
 [3 6]]
Matrix Product:
[[14 32]
 [32 77]]
```

10.5 Some Functions Generating array Objects

```
[20]: x = np.zeros(4)
print(x)

#The type of zeros (e. g. int, int32, int64, float, float32, float64,
↳ etc.)
#can be specified using the dtype argument:

print(x.dtype)

x = x.astype("int")
print(x.dtype)

y = np.zeros((4,3))
print(y)
```

```
z = np.zeros((2,4,3))
print(z)
```

```
[0. 0. 0. 0.]
float64
int64
[[0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]
 [0. 0. 0.]]
[[[0. 0. 0.]
   [0. 0. 0.]
   [0. 0. 0.]
   [0. 0. 0.]]
 [[0. 0. 0.]
   [0. 0. 0.]
   [0. 0. 0.]
   [0. 0. 0.]]]
```

```
[15]: x = list(range(1,11))
print(x)

y = np.array(x).reshape(2,5)
print(y)

# Two ways to use shape
print(y.shape)
np.shape(y)
```

```
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
[[ 1  2  3  4  5]
 [ 6  7  8  9 10]]
(2, 5)
```

```
[15]: (2, 5)
```

```
[33]: shape = np.shape(y)
# Number of rows
n_rows = shape[0]
print("Number of rows:", n_rows)

# Number of columns
n_cols = shape[1]
print("Number of columns:", n_cols)

print(y.shape[0], y.shape[1])
```

```
Number of rows: 2
Number of columns: 5
2 5
```

R is particularly well-suited for statistical and matrix operations, and as such, it often provides more straightforward ways to perform these kinds of tasks. The matrix function in R is a great example of this. R's matrix function directly takes the data (or, the vector x) and the desired dimensions (number of rows and columns) as arguments, making it very concise and easy to use for creating matrices.

This ease of use for matrix and array operations is one of the reasons why R is so popular in data analysis, statistics, and related fields. Python also has powerful capabilities for these tasks, especially with libraries like NumPy, but the syntax can be more verbose compared to R for certain operations.

```
[12]: %%R
y <- matrix(1:10, 2, 5)
print(y)
dim(y)
```

```
      [,1] [,2] [,3] [,4] [,5]
[1,]     1     3     5     7     9
[2,]     2     4     6     8    10
[1] 2 5
```

10.6 Extracting Elements from an Array

Access to the elements of an array is done in the same way as for lists, using indexes. The syntax is as follows: `array[lower:upper:step]` - When `lower` is not specified, the first element (indexed 0) is considered as the value assigned to lower. - When `upper` is not specified, the last element is considered as the value assigned to upper. - When `step` is not specified, a step of 1 is assigned by default.

```
[21]: print(y)
y[1,3]
```

```
[[ 1  2  3  4  5]
 [ 6  7  8  9 10]]
```

```
[21]: 9
```

```
[22]: y[,3]
```

```
Cell In[22], line 1
```

```
y[,3]
```

```
~
```

```
SyntaxError: invalid syntax
```

```
[23]: y[:,3]
```

```
[23]: array([4, 9])
```

```
[27]: y[0,:]
```

```
[27]: array([1, 2, 3, 4, 5])
```

```
[28]: y[0,:-1]
```

```
[28]: array([1, 2, 3, 4])
```

```
[36]: y[0,2:4]
```

```
[36]: array([3, 4])
```

```
[29]: len(y[0,:-1])
```

```
[29]: 4
```

```
[38]: A = np.array(np.arange(16)).reshape((4, 4))  
A
```

```
[38]: array([[ 0,  1,  2,  3],  
           [ 4,  5,  6,  7],  
           [ 8,  9, 10, 11],  
           [12, 13, 14, 15]])
```

```
[37]: A[1,2]
```

```
[37]: 6
```

To select multiple rows at a time, we can pass in a list specifying our selection. For instance, `[1,2]` will retrieve the second and third rows:

```
[39]: A[[1,2]]
```

```
[39]: array([[ 4,  5,  6,  7],  
           [ 8,  9, 10, 11]])
```

To select the first and third columns, we pass in `[0,2]` as the second argument in the square brackets. In this case we need to supply the first argument : which selects all rows.

```
[40]: A[:,[0,2]]
```

```
[40]: array([[ 0,  2],  
           [ 4,  6],  
           [ 8, 10],
```

```
[12, 14]])
```

Now, suppose that we want to select the submatrix made up of the second and fourth rows as well as the first and third columns. This is where indexing gets slightly tricky. It is natural to try to use lists to retrieve the rows and columns:

```
[41]: A[[1,3],[0,2]]
```

```
[41]: array([ 4, 14])
```

When supplied with two indexing lists, the numpy interpretation is that these provide pairs of *i, j* indices for a series of entries. That is why the pair of lists must have the same length. However, that was not our intent, since we are looking for a submatrix. One easy way to do this is as follows. We first create a submatrix by subsetting the rows of *A*, and then on the fly we make a further submatrix by subsetting its columns.

```
[42]: A[[1,3]][:,[0,2]]
```

```
[42]: array([[ 4,  6],
            [12, 14]])
```

An easier way:

```
[43]: idx = np.ix_([1,3],[0,2,3])
      A[idx]
```

```
[43]: array([[ 4,  6,  7],
            [12, 14, 15]])
```

The convenience function `np.ix_()` allows us to extract a submatrix using lists, by creating an intermediate mesh object.

Here, the function `np.all()` has checked whether all entries of an array are `True`. A similar function, `np.any()`, can be used to check whether any entries of an array are `True`.

```
[44]: np.all(A == 0)
```

```
[44]: False
```

```
[45]: np.any(A == 6)
```

```
[45]: True
```

10.7 Missing Values

In Python, especially when working with data science and machine learning, missing values are commonly represented using `None` or `NaN` (Not a Number). The choice between these two depends on the context and the type of data you are dealing with. **None**: This is the Python equivalent of `NULL` in other languages. It's a Python object that often represents the absence of a value. `None` is

used in general Python programming to denote the absence of a value or a null state. **NaN**: This is a special floating-point value defined in the IEEE floating-point standard. It is used to represent missing or undefined numerical data. In Python, NaN is used primarily in numerical arrays and is part of the NumPy library, which is extensively used in data science and machine learning.

When working with pandas, a popular data manipulation library in Python, both **None** and **NaN** can represent missing data. Pandas treats **None** and **NaN** as essentially interchangeable for indicating missing or null values. To facilitate this convention, there are several methods for detecting, removing, and replacing null values in pandas

In NumPy, which is a core library for numerical computing in Python, missing values are typically represented using `numpy.nan` for floating-point data. Unlike Python's general-purpose **None**, which can be used with any data type, `numpy.nan` is specifically a floating-point value and fits naturally into arrays of floats.

```
[47]: a = np.array([1.0, 2.0, NaN, 4.0])
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[47], line 1
----> 1 a = np.array([1.0, 2.0, NaN, 4.0])

NameError: name 'NaN' is not defined
```

```
[57]: a = np.array([1.0, 2.0, None, 4.0])
a
```

```
[57]: array([1.0, 2.0, None, 4.0], dtype=object)
```

```
[60]: b = np.array([1.0, 2.0, np.nan, 4.0])
b
```

```
[60]: array([ 1.,  2., nan,  4.])
```

```
[51]: np.isnan(a)
```

```
[51]: array([False, False,  True, False])
```

```
[61]: np.isnan(b)
```

```
[61]: array([False, False,  True, False])
```

```
[52]: np.isnan(a).any()
```

```
[52]: True
```

```
[53]: # Calculate the mean, ignoring NaN values
np.nanmean(a)
```

```
[53]: 2.3333333333333335
```

```
[62]: np.nanmean(b)
```

```
[62]: 2.3333333333333335
```

```
[54]: # Example arrays
b = np.array([10, 20, 30, 40])
c = np.array([2, 0, 5, 0])

d = np.divide(b, c)

print(d)
```

```
[ 5. inf  6. inf]
```

```
/var/folders/wt/4xtk6v051vd349k3wktfld480000gn/T/ipykernel_55077/1948954575.py:6
: RuntimeWarning: divide by zero encountered in divide
d = np.divide(b, c)
```

```
[56]: # Safe division
d = np.divide(b, c, where=c!=0)

print(d)
```

```
[ 5. -0.  6. -0.]
```

```
[63]: import pandas as pd

# Create a DataFrame with missing values
# Will See Pandas in CH.10 Below
df = pd.DataFrame({'A': [1, 2, None], 'B': [4, None, 6], 'C': [7, 8, np.nan]})

# Output the DataFrame
print(df)
```

	A	B	C
0	1.0	4.0	7.0
1	2.0	NaN	8.0
2	NaN	6.0	NaN

10.8 Conditional indexing

Conditional indexing in Python, especially with NumPy arrays and Pandas DataFrames, is a powerful feature that allows you to select elements based on some condition. Let's go through examples for both NumPy arrays and Pandas DataFrames:

```
[41]: import numpy as np

# Creating a NumPy array
arr = np.array([1, 2, 3, 4, 5, 6, 7, 8, 9, 10])

# Condition: Select elements greater than 5
selected_elements = arr[arr > 5]
print(selected_elements)

# Creating a 2D NumPy array (matrix)
matrix = np.array([[1, 2, 3],
                   [4, 5, 6],
                   [7, 8, 9]])

# Condition: Select elements greater than 5
selected_elements = matrix[matrix > 5]
print(selected_elements)

matrix[matrix > 5]
```

```
[ 6  7  8  9 10]
[6 7 8 9]
```

[41]: False

```
[43]: # Condition: Find indices of elements greater than 5
indices = np.where(matrix > 5)

print(indices)
print(indices[0])
print(indices[1])
```

```
(array([1, 2, 2, 2]), array([2, 0, 1, 2]))
[1 2 2 2]
[2 0 1 2]
```

```
[44]: # Using indices to extract the values from the matrix
selected_values = matrix[indices]

print(selected_values)
```

```
[6 7 8 9]
```

```
[37]: import pandas as pd

# Creating a DataFrame
df = pd.DataFrame({
    'A': [1, 2, 3, 4, 5],
```



```

    'B': [5, 4, 3, 2, 1],
    'C': [2, 3, 4, 5, 6]
})

# Condition: Select rows where column 'A' is greater than 2
selected_rows = df[df['A'] > 2]
print(selected_rows)

# Condition: Select rows where 'A' is greater than 2 and 'B' is less than 5
selected_rows = df[(df['A'] > 2) & (df['B'] < 5)]
print(selected_rows)

```

```

    A  B  C
2  3  3  4
3  4  2  5
4  5  1  6
    A  B  C
2  3  3  4
3  4  2  5
4  5  1  6

```

```

[49]: %%R

X <- matrix(1:10, 2, 5)
print(X[X>5])

# or

ind <- which(X>5, arr.ind = TRUE)
print(ind)

print(X[ind])

```

```

[1]  6  7  8  9 10
    row col
[1,]   2   3
[2,]   1   4
[3,]   2   4
[4,]   1   5
[5,]   2   5
[1]  6  7  8  9 10

```

```

[53]: %%R

# Creating a DataFrame
df = data.frame(A = 1:5,
                B = 5:1,

```

```

C = 2:6)

# Condition: Select rows where column 'A' is greater than 2
selected_rows = df[df$A > 2, ]
print(selected_rows)

# Condition: Select rows where 'A' is greater than 2 and 'B' is less than 5
selected_rows = df[df$A > 2 & df$B < 5, ]
print(selected_rows)

# which()

ind <- which(df$A > 2 & df$B < 5)
selected_rows = df[ind, ]
print(selected_rows)

```

```

  A B C
3 3 3 4
4 4 2 5
5 5 1 6
  A B C
3 3 3 4
4 4 2 5
5 5 1 6
  A B C
3 3 3 4
4 4 2 5
5 5 1 6

```

R has a reputation for being particularly concise and straightforward for certain types of data operations, especially those involving matrices and data frames. This is largely because R is specifically designed for statistical computing and data analysis, so operations like matrix manipulation are very streamlined.

In R, functions like `rbind()` and `cbind()` are directly tailored for adding rows and columns to matrices, making the syntax very clean and intuitive for these operations. This focus on data-centric tasks is one of the strengths of R and a reason why it's so popular in the statistics and data science communities.

Python, being a more general-purpose language, can require more verbose code for the same operations, particularly when using libraries like NumPy for numerical computing. However, Python's broad applicability across different domains, from web development to machine learning, makes it a versatile choice.

Both languages have their strengths, and the choice often depends on the specific needs and context of the task at hand. For pure data analysis and statistical work, R might be more convenient, while Python offers greater flexibility for a wide range of applications.

```
[54]: %%R

mat <- matrix(1:9, 3, 3)
r <- c(1,1,-5)
co <- c(19, 0.57, 22, 42)

mat1 <- rbind(mat, r)
mat2 <- cbind(mat1, co)
mat2
```

```
      co
1 4  7 19.00
2 5  8  0.57
3 6  9 22.00
r 1 1 -5 42.00
```

```
[2]: import numpy as np

# Create a 3x3 matrix
mat = np.array([[1, 2, 3],
                [4, 5, 6],
                [7, 8, 9]])

# New row and column to be added
r = np.array([1, 1, -5])
co = np.array([19, 0.57, 22, 42])

# Add the new row to the bottom of the matrix
mat1 = np.vstack([mat, r])

# Add the new column to the right of the matrix
# First, reshape 'co' to be a 4x1 column vector
co = co.reshape(-1, 1) # '-1' lets NumPy automatically calculate the size
mat2 = np.hstack([mat1, co])

print("Original Matrix:\n", mat)
print("Matrix with New Row:\n", mat1)
print("Matrix with New Row and Column:\n", mat2)
```

Original Matrix:

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

Matrix with New Row:

```
[[ 1  2  3]
 [ 4  5  6]
 [ 7  8  9]
 [ 1  1 -5]]
```

Matrix with New Row and Column:

```
[[ 1.   2.   3.  19. ]
 [ 4.   5.   6.   0.57]
 [ 7.   8.   9.  22. ]
 [ 1.   1.  -5.  42. ]]
```

In Python, the `append` method for lists and the `numpy.append` function for NumPy arrays work differently than the `vstack` and `hstack` functions we used above. Let's clarify the differences:

List `append` Method in Python: - The `append` method is used with Python lists to add a single element to the end of the list. - It does not return a new list; rather, it modifies the existing list in place. - It cannot be used directly for adding a row or a column to a 2D list or a NumPy array in the way `vstack` or `hstack` do.

```
[ ]: my_list = [1, 2, 3]
      my_list.append(4)  # Adds 4 to the end of my_list
      print(my_list)    # Output: [1, 2, 3, 4]
```

NumPy `append` Function: - `numpy.append` adds values to the end of an array. - Unlike list `append`, it returns a new array and does not modify the original array in place. - It flattens the array by default if you don't specify the axis, which means it will return a 1D array even if you append a row or a column to a 2D array. - To use it for adding a row or a column without flattening, you need to specify the axis parameter.

```
[58]: mat = np.array([[1, 2, 3], [4, 5, 6]])
      new_row = np.array([7, 8, 9])

      # Appending a new row without flattening
      mat_with_row = np.append(mat, [new_row], axis=0)
      print(mat_with_row)

      new_row = np.array([7, 8, 9])

      # Appending a new column without flattening
      new_col = np.array([7, 8])
      mat_with_col = np.append(mat, [new_col], axis=1)
      print(mat_with_col)
```

```
[[1 2 3]
 [4 5 6]
 [7 8 9]]
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[58], line 12
     10 # Appending a new column without flattening
     11 new_col = np.array([7, 8])
--> 12 mat_with_col = np.append(mat, [new_col], axis=1)
     13 print(mat_with_col)
```

```

File <__array_function__ internals>:180, in append(*args, **kwargs)

File ~/anaconda3/lib/python3.10/site-packages/numpy/lib/function_base.py:5444, in
    ↪in append(arr, values, axis)
    5442     values = ravel(values)
    5443     axis = arr.ndim-1
-> 5444 return concatenate((arr, values), axis=axis)

File <__array_function__ internals>:180, in concatenate(*args, **kwargs)

ValueError: all the input array dimensions for the concatenation axis must match
    ↪exactly, but along dimension 0, the array at index 0 has size 2 and the array
    ↪at index 1 has size 1

```

For column-wise appending (along `axis=1`), as in appending `new_col` to `mat`, you often need to reshape the array to match the number of rows in `mat`, especially if `new_col` starts as a 1D array. But for row-wise appending (along `axis=0`), as long as the number of elements in the row being appended matches the number of columns in the matrix, additional reshaping isn't necessary.

Also note that, in our row-wise appending above, when appending `new_row` to `mat` using `np.append`, the square brackets `[new_row]` are used to ensure that `new_row` is treated as a two-dimensional array with one row, rather than as a one-dimensional array. This is important for aligning the dimensions correctly when appending along `axis=0` (the row axis).

```

[59]: new_col = np.array([7, 8])
      mat_with_col = np.append(mat, new_col.reshape(2, 1), axis=1)
      print(mat_with_col)

```

```

[[1 2 3 7]
 [4 5 6 8]]

```

10.9 Sequences and Slice Notations

The function `np.linspace()` can be used to create a sequence of numbers.

```

[32]: seq1 = np.linspace(0, 10, 11)
      seq1

```

```

[32]: array([ 0.,  1.,  2.,  3.,  4.,  5.,  6.,  7.,  8.,  9., 10.])

```

The function `np.arange()` returns a sequence of numbers spaced out by `step`. If `step` is not specified, then a default value of 1 is used. Let's create a sequence that starts at 0 and ends at 10.

```

[33]: seq2 = np.arange(0,10)
      seq2

```

```

[33]: array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])

```

Why isn't 10 output above? This has to do with slice notation in Python. Slice notation is used to index sequences such as lists, tuples and arrays. Suppose we want to retrieve the fourth through sixth (inclusive) entries of a string. We obtain a slice of the string using the indexing notation `[3:6]`

```
[34]: "hello world"[3:6]
```

```
[34]: 'lo '
```

In the code block above, the notation `3:6` is shorthand for `slice(3,6)` when used inside `[]`.

```
[35]: "hello world"[slice(3,6)]
```

```
[35]: 'lo '
```

You might have expected `slice(3,6)` to output the fourth through seventh characters in the text string (recalling that Python begins its indexing at zero), but instead it output the fourth through sixth. This also explains why the earlier `np.arange(0, 10)` command output only the integers from 0 to 9. See the documentation [slice?](#) for useful options in creating slices

```
[36]: A = np.array(np.arange(16)).reshape((4, 4))
      A
```

```
[36]: array([[ 0,  1,  2,  3],
             [ 4,  5,  6,  7],
             [ 8,  9, 10, 11],
             [12, 13, 14, 15]])
```


11 Graphics

In Python, common practice is to use the library `matplotlib` for graphics. However, since Python was not written with data analysis in mind, the notion of plotting is not intrinsic to the language. We will use the `subplots()` function from `matplotlib.pyplot` to create a figure and the axes onto which we plot our data. For many more examples of how to make plots in Python, readers are encouraged to visit matplotlib.org/stable/gallery/.


```
[5]: from matplotlib.pyplot import subplots

     # Create a random number generator instance
     rng = np.random.default_rng()


     fig, ax = subplots(figsize=(5, 5))
     x = rng.standard_normal(100)
     y = rng.standard_normal(100)
     ax.plot(x, y)
```



```
[6]: fig, ax = subplots(figsize=(5, 5))  
     ax.plot(x, y, 'o')
```



```
[8]: fig, ax = subplots(figsize=(5, 5))
      ax.scatter(x, y, marker='o')
      ax.set_xlabel("this is the x-axis")
      ax.set_ylabel("this is the y-axis")
      ax.set_title("Plot of X vs Y");
```

```
[27]: import matplotlib.pyplot as plt

# Clear any existing plots
plt.clf()

# Create a random number generator instance
rng = np.random.default_rng()

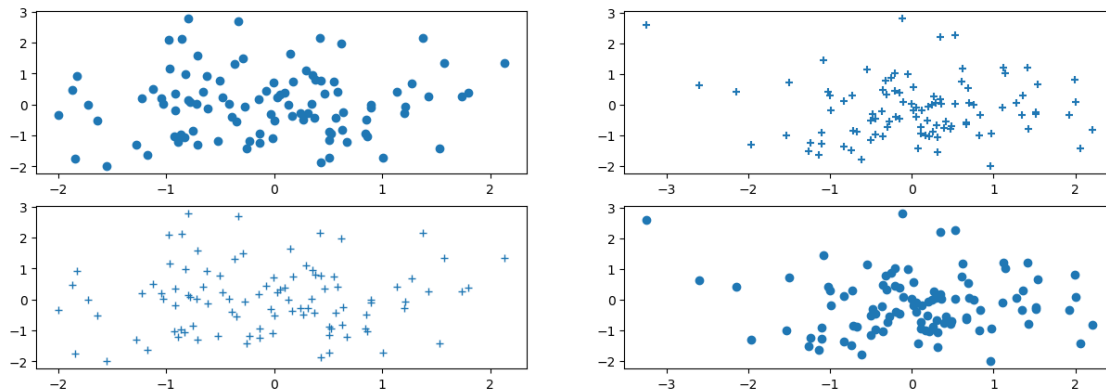
# Generate random data
x = rng.standard_normal(100)
y = rng.standard_normal(100)
z = rng.standard_normal(100)
w = rng.standard_normal(100)

# Create a 2x2 grid of subplots
fig, axes =.subplots(nrows=2, ncols=2, figsize=(15, 5))

# Plot in the first subplot
axes[0,0].plot(x, y, 'o')
```

```
axes[1,0].plot(x, y, '+')
# Scatter plot in the second subplot
axes[0,1].scatter(z, w, marker='+')
axes[1,1].scatter(z, w, marker='o')
plt.show()
```

<Figure size 640x480 with 0 Axes>



12 Pandas

Data sets often contain different types of data, and may have names associated with the rows or columns. For these reasons, they typically are best accommodated using a data frame. We can think of a data frame as a sequence of arrays of identical length; these are the columns. Entries in the different arrays can be combined to form a row. The pandas library can be used to create and work with data frame objects.

Reading in a Data Set The first step of most analyses involves importing a data set into Python. Before attempting to load a data set, we must make sure that Python knows where to find the file containing it. If the file is in the same location as this notebook file, then we are all set. Otherwise, the command `os.chdir()` can be used to change directory. (You will need to call `import os` before calling `os.chdir()`.) We will begin by reading in `Auto.csv`. This is a comma-separated file, and can be read in using `pd.read_csv()`:

```
[3]: import numpy as np
import pandas as pd
Auto = pd.read_csv("Auto.csv")
Auto
```

```
[3]:      mpg  cylinders  displacement  horsepower  weight  acceleration  year  \
0    18.0         8         307.0         130    3504         12.0    70
1    15.0         8         350.0         165    3693         11.5    70
2    18.0         8         318.0         150    3436         11.0    70
```

3	16.0	8	304.0	150	3433	12.0	70
4	17.0	8	302.0	140	3449	10.5	70
..
392	27.0	4	140.0	86	2790	15.6	82
393	44.0	4	97.0	52	2130	24.6	82
394	32.0	4	135.0	84	2295	11.6	82
395	28.0	4	120.0	79	2625	18.6	82
396	31.0	4	119.0	82	2720	19.4	82

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino
..
392	1	ford mustang gl
393	2	vw pickup
394	1	dodge rampage
395	1	ford ranger
396	1	chevy s-10

[397 rows x 9 columns]

```
[2]: Auto['horsepower']
```

```
[2]: 0      130
      1      165
      2      150
      3      150
      4      140
      ...
      392     86
      393     52
      394     84
      395     79
      396     82
```

Name: horsepower, Length: 397, dtype: object

```
[5]: np.unique(Auto['horsepower'])
```

```
[5]: array(['100', '102', '103', '105', '107', '108', '110', '112', '113',
          '115', '116', '120', '122', '125', '129', '130', '132', '133',
          '135', '137', '138', '139', '140', '142', '145', '148', '149',
          '150', '152', '153', '155', '158', '160', '165', '167', '170',
          '175', '180', '190', '193', '198', '200', '208', '210', '215',
          '220', '225', '230', '46', '48', '49', '52', '53', '54', '58',
```

```
'60', '61', '62', '63', '64', '65', '66', '67', '68', '69', '70',
'71', '72', '74', '75', '76', '77', '78', '79', '80', '81', '82',
'83', '84', '85', '86', '87', '88', '89', '90', '91', '92', '93',
'94', '95', '96', '97', '98', '?'], dtype=object)
```

We see the value ? is being used to encode missing values.

```
[6]: Auto = pd.read_csv("Auto.csv",
                        na_values = ["?"])
np.unique(Auto['horsepower'])
```

```
[6]: array([ 46.,  48.,  49.,  52.,  53.,  54.,  58.,  60.,  61.,  62.,  63.,
          64.,  65.,  66.,  67.,  68.,  69.,  70.,  71.,  72.,  74.,  75.,
          76.,  77.,  78.,  79.,  80.,  81.,  82.,  83.,  84.,  85.,  86.,
          87.,  88.,  89.,  90.,  91.,  92.,  93.,  94.,  95.,  96.,  97.,
          98., 100., 102., 103., 105., 107., 108., 110., 112., 113., 115.,
         116., 120., 122., 125., 129., 130., 132., 133., 135., 137., 138.,
         139., 140., 142., 145., 148., 149., 150., 152., 153., 155., 158.,
         160., 165., 167., 170., 175., 180., 190., 193., 198., 200., 208.,
         210., 215., 220., 225., 230.,  nan])
```

```
[7]: Auto['horsepower'].sum()
```

```
[7]: 40952.0
```

```
[33]: Auto.shape
```

```
[33]: (397, 9)
```

```
[35]: Auto.isna().any(axis=1).sum()
```

```
[35]: 5
```

There are various ways to deal with missing data. In this case, since only five of the rows contain missing observations, we choose to use the `Auto.dropna()` method to simply remove these rows.

```
[4]: Auton = Auto.dropna()
Auton.shape
```

```
[4]: (397, 9)
```

```
[40]: Auton.columns
```

```
[40]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
          'acceleration', 'year', 'origin', 'name'],
          dtype='object')
```

Accessing the rows and columns of a data frame is similar, but not identical, to accessing the rows and columns of an array. Recall that the first argument to the `[]` method is always applied to the rows of the array. Similarly, passing in a slice to the `[]` method creates a data frame whose rows are determined by the slice:

```
[38]: Auton[:3]
```

```
[38]:      mpg  cylinders  displacement  horsepower  weight  acceleration  year  \
0   18.0           8         307.0        130.0    3504           12.0    70
1   15.0           8         350.0        165.0    3693           11.5    70
2   18.0           8         318.0        150.0    3436           11.0    70

      origin          name
0         1  chevrolet chevelle malibu
1         1         buick skylark 320
2         1    plymouth satellite
```

```
[45]: # Conditional Indexing
idx_80 = Auton['weight'] > 5000
Auton[idx_80]
```

```
[45]:      mpg  cylinders  displacement  horsepower  weight  acceleration  year  \
44  13.0           8         400.0        175.0    5140           12.0    71

      origin          name
44         1  pontiac safari (sw)
```

If we pass in a list of strings to the `[]` method, we obtain a data frame containing the corresponding set of columns.

```
[46]: Auton[['mpg', 'horsepower']]
```

```
[46]:      mpg  horsepower
0   18.0        130.0
1   15.0        165.0
2   18.0        150.0
3   16.0        150.0
4   17.0        140.0
..   ...          ...
392  27.0         86.0
393  44.0         52.0
394  32.0         84.0
395  28.0         79.0
396  31.0         82.0
```

```
[392 rows x 2 columns]
```

Using index for selection In database systems, an index is a data structure that improves the speed of data retrieval operations on a database table. Indexes are used to quickly locate data without having to search every row in a database table every time a database table is accessed. When you apply a filter or a query on a column that has an index, the database management system can use the index to find rows faster. This is analogous to an index in a book - instead of reading every page, you refer to the index to find the pages you need. For example, if you have an indexed column `employee_id`, and you query for `employee_id = 12345`, the database can use the index to quickly locate this record.

In data analysis tools like Pandas in Python, indexing can also help in faster data retrieval and filtering. Pandas allows you to set indexes on DataFrames for similar purposes.

When we loaded our data frame, the rows are labeled using integers 0 to 396.

```
[47]: Auton.index
```

```
[47]: Int64Index([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9,
...
               387, 388, 389, 390, 391, 392, 393, 394, 395, 396],
              dtype='int64', length=392)
```

```
[5]: Auto_re = Auton.set_index('name')
      Auto_re
```

```
[5]:
```

	mpg	cylinders	displacement	horsepower	weight	\
name						
chevrolet chevelle malibu	18.0	8	307.0	130	3504	
buick skylark 320	15.0	8	350.0	165	3693	
plymouth satellite	18.0	8	318.0	150	3436	
amc rebel sst	16.0	8	304.0	150	3433	
ford torino	17.0	8	302.0	140	3449	
...	
ford mustang gl	27.0	4	140.0	86	2790	
vw pickup	44.0	4	97.0	52	2130	
dodge rampage	32.0	4	135.0	84	2295	
ford ranger	28.0	4	120.0	79	2625	
chevy s-10	31.0	4	119.0	82	2720	

	acceleration	year	origin
name			
chevrolet chevelle malibu	12.0	70	1
buick skylark 320	11.5	70	1
plymouth satellite	11.0	70	1
amc rebel sst	12.0	70	1
ford torino	10.5	70	1
...
ford mustang gl	15.6	82	1
vw pickup	24.6	82	2

dodge rampage	11.6	82	1
ford ranger	18.6	82	1
chevy s-10	19.4	82	1

[397 rows x 8 columns]

```
[49]: # We see that the column 'name' is no longer there.
Auto_re.columns
```

```
[49]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
          'acceleration', 'year', 'origin'],
          dtype='object')
```

Index entries need not be unique: there are several cars in the data frame named `ford galaxie 500`.

```
[56]: Auto_re.loc['ford galaxie 500', ['mpg', 'origin']]
```

```
[56]:
```

	mpg	origin
name		
ford galaxie 500	15.0	1
ford galaxie 500	14.0	1
ford galaxie 500	14.0	1

Now, we can access rows of the data frame by name using the `loc[]` method of `Auto`:

```
[50]: rows = ['amc rebel sst', 'ford torino']
Auto_re.loc[rows]
```

```
[50]:
```

	mpg	cylinders	displacement	horsepower	weight	\
name						
amc rebel sst	16.0	8	304.0	150.0	3433	
ford torino	17.0	8	302.0	140.0	3449	

	acceleration	year	origin
name			
amc rebel sst	12.0	70	1
ford torino	10.5	70	1

```
[53]: Auto_re.iloc[[3,4]]
```

```
[53]:
```

	mpg	cylinders	displacement	horsepower	weight	\
name						
amc rebel sst	16.0	8	304.0	150.0	3433	
ford torino	17.0	8	302.0	140.0	3449	

	acceleration	year	origin
name			

amc rebel sst	12.0	70	1
ford torino	10.5	70	1

`loc[]` is label-based; it uses the names or labels of rows/columns. `iloc[]` is integer position-based; it uses the numerical positions in the DataFrame, starting from 0. The same results from both methods indicate a match between the label positions and their numerical positions in this particular instance of your DataFrame.

```
[54]: # We can also use it to retrieve the 1st, 3rd and 4th columns
Auto_re.iloc[:, [0,2,3]]
```

```
[54]:
```

	mpg	displacement	horsepower
name			
chevrolet chevelle malibu	18.0	307.0	130.0
buick skylark 320	15.0	350.0	165.0
plymouth satellite	18.0	318.0	150.0
amc rebel sst	16.0	304.0	150.0
ford torino	17.0	302.0	140.0
...
ford mustang gl	27.0	140.0	86.0
vw pickup	44.0	97.0	52.0
dodge rampage	32.0	135.0	84.0
ford ranger	28.0	120.0	79.0
chevy s-10	31.0	119.0	82.0

[392 rows x 3 columns]

We can extract the 4th and 5th rows, as well as the 1st, 3rd and 4th columns, using a single call to `iloc[]`:

```
[55]: Auto_re.iloc[[3,4], [0,2,3]]
```

```
[55]:
```

	mpg	displacement	horsepower
name			
amc rebel sst	16.0	304.0	150.0
ford torino	17.0	302.0	140.0

Suppose that we want to create a data frame consisting of the year and origin of the subset of cars with weight greater than 5000

```
[60]: idx_80 = Auto_re['weight'] > 5000
Auto_re.loc[idx_80, ['weight', 'origin', 'year']]
```

```
[60]:
```

	weight	origin	year
name			
pontiac safari (sw)	5140	1	71

To do this more concisely, we can use an anonymous function called a `lambda`:


```
[61]: Auto_re.loc[lambda df: df['weight'] > 5000, ['weight', 'origin', 'year']]
```

```
[61]:
```

	weight	origin	year
name			
pontiac safari (sw)	5140	1	71

As another example of using a `lambda`, suppose that we want all cars built after 1981 that achieve greater than 30 miles per gallon:

```
[5]: Auto_re.loc[lambda df: (df['year'] > 81) & (df['mpg'] > 30),
                ['weight', 'origin']]
```

```
[5]:
```

	weight	origin
name		
chevrolet cavalier 2-door	2395	1
pontiac j2000 se hatchback	2575	1
volkswagen rabbit l	1980	2
mazda glc custom l	2025	3
mazda glc custom	1970	3
plymouth horizon miser	2125	1
mercury lynx l	2125	1
nissan stanza xe	2160	3
honda accord	2205	3
toyota corolla	2245	3
honda civic	1965	3
honda civic (auto)	1965	3
datsum 310 gx	1995	3
oldsmobile cutlass ciera (diesel)	3015	1
toyota celica gt	2665	3
dodge charger 2.2	2370	1
vw pickup	2130	2
dodge rampage	2295	1
chevy s-10	2720	1

As another example, suppose that we want to retrieve all Ford and Datsun cars with displacement less than 80. We check whether each `name` entry contains either the string `ford` or `datsum` using the `str.contains()` method of the index attribute of the dataframe:

```
[70]: Auto_re.loc[lambda df: (df['displacement'] < 80)
                    & (df.index.str.contains('ford')
                     | df.index.str.contains('datsum')),
                ['weight', 'origin']]
```

```
[70]:
```

	weight	origin
name		
datsum 1200	1613	3

```
datsum b210    1950      3
```

To replicate this type of data filtering in R, we can use a combination of base R functions and, optionally, functions from the `dplyr` package, which is part of the `tidyverse` and provides a more intuitive syntax for data manipulation. Indexing doesn't work in R as it works in Python. Here's how you can achieve the same result:

```
[17]: %%R -i Auton

# Filter for Ford and Datsun cars with displacement < 80
filtered_cars <- Auton[with(Auton, displacement < 80 &
                           (grepl("ford", name, ignore.case = TRUE) |
                            grepl("datsum", name, ignore.case = TRUE))),
                      c("name", "weight", "origin")]

filtered_cars
```

```
      name weight origin
54  datsum 1200   1613     3
129 datsum b210   1950     3
```

If you're using the `dplyr` package, you can achieve the same with a more readable syntax:

```
[25]: %%R -i Auton

library(dplyr)
library(tibble)

# Filter using dplyr
fil_cars <- Auton %>%
  filter(displacement < 80,
         grepl("ford", name, ignore.case = TRUE) |
         grepl("datsum", name, ignore.case = TRUE)) %>%
  select(name, weight, origin)


fil_cars
```

```
      name weight origin
54  datsum 1200   1613     3
129 datsum b210   1950     3
```

13 Plots and Numerical Summaries

```
[38]: from matplotlib.pyplot import subplots

fig, ax = subplots(figsize=(8, 8))
ax.plot(Auto['horsepower'], Auto['mpg'], 'o');
```




```
[40]: # Here we use Auto.plot() so the variables can be accessed by
      # their names

      ax = Auto.plot.scatter('horsepower', 'mpg')
      ax.set_title('Horsepower vs. MPG')

      # If you want to save
      # fig = ax.figure
      # fig.savefig("name_of_the_figure.png")
```

```
[40]: Text(0.5, 1.0, 'Horsepower vs. MPG')
```



Consider the `cylinders` variable. Typing in `Auto.cylinders.dtype` reveals that it is being treated as a quantitative variable. However, since there is only a small number of possible values for this variable, we may wish to treat it as qualitative. Below, we replace the `cylinders` column with a categorical version of `Auto.cylinders`. The function `pd.Series()` owes its name to the fact that pandas is often used in time series applications.

```
[41]: Auto.cylinders.dtype
```


```
[41]: dtype('int64')
```

```
[42]: Auto.cylinders = pd.Series(Auto.cylinders, dtype='category')
Auto.cylinders.dtype
```

```
[42]: CategoricalDtype(categories=[3, 4, 5, 6, 8], ordered=False)
```

Now that `cylinders` is qualitative, we can display it using the `boxplot()` or the `hist()` method.

```
[43]: fig, ax = subplots(figsize=(8, 8))
Auto.hist('mpg', ax=ax);
```



The `describe()` method produces a numerical summary of each column in a data frame.

```
[44]: Auto[['mpg', 'weight']].describe()
```

```
[44]:
```

	mpg	weight
count	397.000000	397.000000
mean	23.515869	2970.261965
std	7.825804	847.904119
min	9.000000	1613.000000
25%	17.500000	2223.000000
50%	23.000000	2800.000000
75%	29.000000	3609.000000
max	46.600000	5140.000000

14 Linear Regression

```
[ ]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
```

14.1 TBC:-)